

AI Work Assistant Handbook

Build repeatable AI workflows: assign clearly, verify output, refine results, and save what works.

Handbook Information

Field	Details
Edition	Free Public Edition v1.0
Last Updated	2026-05-08
Author	Farceur Liu
Audience	First-time AI users in professional settings; individuals, small teams, and managers looking to integrate AI into daily work
Core Theme	Start with one low-risk task, then build a personal system for assigning work to AI, verifying output, refining results, and saving what works
Suggested Use	Self-study, team onboarding, workflow documentation, skill design
Not Suitable For	Replacing legal, financial, security, medical, or HR decisions

About This Handbook

Most people who start using AI ask the same first question: "Is there a better tool?"

But the real obstacle is rarely the tool. It is usually the process: not knowing how to delegate clearly, how to provide the right context, how to stop AI from guessing, how to review the first draft, or how to save a useful method so it can be reused.

The core argument of this handbook is simple: you do not need many AI tools. One or two mainstream tools are almost always enough. The real gap is not the tool list - it is the working method.

This handbook starts with one low-risk task and walks you through turning AI from a chat partner into a work assistant you can actually delegate to. You will learn how to choose tasks, write clear assignments, read for errors, ask for revisions, save templates, and decide when a workflow is worth packaging as a skill.

A Note on the Original Title

The original Chinese version of this handbook is titled: "從零開始養成我的 AI 管家." A rough translation would be "Building My AI Butler from Scratch." The word "butler" in that title evokes the idea of a trusted assistant who learns your preferences, handles repeatable tasks, and follows standing instructions - but never makes final decisions on your behalf.

This English edition uses "AI work assistant" instead. The idea is the same: you are not trying to find the perfect tool or automate everything at once. You are building a working relationship with AI - teaching it your context, setting clear boundaries, and gradually turning reliable methods into reusable systems.

Public Disclaimer and Permissions

This handbook provides general guidance on AI work methods and safety principles. It does not constitute legal, financial, security, medical, or HR advice. When working with company data, customer data, personal information, financial transactions, contracts, account permissions, system operations, script execution, or file deletion, you must follow your organization's policies, authorization scope, and IT or security guidelines.

Do not use examples from this handbook to bypass organizational policies, install unknown tools, run scripts you do not understand, or process unauthorized data. Tool names, interfaces, and features change over time - refer to official product documentation and your organization's policies for current guidance.

The product names ChatGPT, Claude, Gemini, Codex, Claude Code, Gemini CLI, and others mentioned in this handbook are the property of their respective owners. This handbook does not represent any affiliation, authorization, endorsement, or guarantee relationship with those tools, platforms, or companies.

Permissions: This handbook may be freely read and shared via its original link. Without explicit written permission from the author, the full text or substantially adapted content may not be resold as commercial courses, paid training materials, publications, consulting deliverables, or enterprise training content. When quoting excerpts, please include the handbook title and source link:

<https://farceurliu.github.io/ai-butler-handbook/>

Before You Begin: Organizational Policies vs. This Handbook

This handbook is a public self-study resource. Its focus is teaching you how to turn AI from a chat tool into a delegatable, verifiable, and accumulating work assistant.

It is not a formal AI usage policy for any company, school, client, or organization. It does not define your approved tools, data classification rules, account permissions, security reporting procedures, legal processes, or financial rules.

Three principles to keep in mind:

1. Methods can be self-studied; data boundaries depend on your organization's rules.
2. Tools can be practiced; formal data and system permissions depend on your authorization.
3. AI can help organize and draft content, but it cannot replace formal decisions, commitments, approvals, or reporting workflows.

If anything in this handbook conflicts with your organization's formal policies, follow your organization's policies.

Table of Contents

- Handbook Information
- About This Handbook
- Public Disclaimer and Permissions

Main Chapters:

- 0. How to Use This Handbook
- 1. What an AI Work Assistant Is, and Is Not
- 2. Your First Low-Risk Assignment
- 3. Beginner No-Go Zones
- 4. Tools, Chat, Workspace AI, CLI Agents, and Skills
- 5. Choosing the Right Tasks
- 6. The Assignment Formula
- 7. Practice Loops
- 8. Everyday Workflow Scenarios
- 9. Diagnosing Bad Assignments
- 10. Verification: Do Not Accept Output Blindly
- 11. Refinement: Make the Second Version Useful
- 12. Safety Boundaries
- 13. Assignment Quality Levels
- 14. Confidence Labels
- 15. Pre-Send Checklist for External Replies
- 16. Template Library
- 17. Red-Flag Rules
- 18. Turning Reusable Workflows into Skills
- 19. Team Adoption and Knowledge Capture
- 20. One-Page Quick Reference
- 21. Closing

Appendices:

- Appendix A: High-Risk Customer Message Triage
- Appendix B: Expense Reimbursement Completeness Check
- Appendix C: Email Search and Attachment Safety Check
- Appendix D: Safe Computer Storage Cleanup

0. How to Use This Handbook

What you will learn in this chapter:

- Understand the recommended reading order and how to navigate this book.
- Know that building an AI work assistant is not a one-time setup — it is a repeating cycle of assigning, reviewing, refining, and saving.
- Establish one core principle before you start: fewer tools used deeply will serve you better than many tools used shallowly.

This book is about building your own AI work assistant from scratch.

"Building" does not mean giving the AI a personality or a name. It means gradually training a general-purpose AI — one that currently just chats — into an assistant that understands your working style. The process looks like this:

1. Give it one low-risk task to start.
2. Review its first attempt.
3. Point out what is wrong and ask for a revised version.
4. Save the method that worked as a reusable template.
5. Repeat until a stable workflow is ready to become a skill.

This book is not meant to be read cover to cover in one sitting. Think of it as a training guide: the early chapters let you start today; the later chapters help you build out a genuine work partner over time.

0.1 If this is your first time using AI as a work tool

Read in this order:

1. **Chapter 1** — Understand what an AI assistant is responsible for (and what it is not)
2. **Chapter 2** — Walk through your first low-risk task
3. **Chapter 6** — Learn the assignment formula
4. **Chapter 7** — Complete the three beginner practice loops
5. **Chapter 10** — Use the four-question review checklist

Your goal the first time is not a perfect output. It is completing one full cycle:

Assign once → Review the first version → Identify the problem → Ask for a revision → Get a more usable second version.

Once you have done that loop even once, you have the foundation for everything else in this book.

0.2 If you just want templates

Go directly to Chapters 8, 16, and 20. But do not copy and paste without adapting. At minimum, fill in three things before you use any template:

- The specific goal of this task
- The data and context AI can refer to
- What AI must not assume, commit to, or modify

A template without these three additions will produce output that looks complete but is not ready to use.

0.3 If your work involves customers, money, personal data, or external-facing content

Read these chapters before you start assigning that type of task:

- **Chapter 3** — Beginner no-go zones
- **Chapter 12** — Safety boundaries
- **Chapter 15** — Pre-send checklist for external replies
- **Chapter 17** — Red-light rules

The core principle for sensitive work is:

AI can help you draft, organize, and check — but it cannot make final commitments on behalf of your organization or replace a human judgment call.

This does not mean AI is useless for these tasks. It means you keep AI in a support role: helping you prepare, not representing you.

0.4 Reader path table

You do not need to read this book from front to back. Find your situation below and follow the matching path.

Your situation	Recommended path	What you should be able to do after
First-time AI user	Chapters 1, 2, 6, 7, 10	Complete a low-risk task and ask for a revised version
Looking for templates or prompts	Chapters 8, 16, 20	Find a reusable template and add your own data, constraints, and verification method
Want to use Codex, Claude Code, or Gemini CLI	Chapters 4, 8.13, 8.15, 18	Distinguish Chat from workspace AI; know which tasks belong in each
Want to create skills	Complete a Chapter 7 or 8 exercise first, then read Chapter 18 and the appendices	Build a skill from a workflow that already runs, not from scratch
Manager or team lead introducing AI to a team	Chapters 3, 12, 17, 19, 20	Establish a low-risk rollout rhythm, shared templates, and human approval checkpoints

If you are new to this, do not jump straight to skills or CLI tools. Build the "assign, review, revise" habit first. The more advanced tools only become useful once that foundation is solid.

0.5 Seven-day practice schedule

If you want a structured first week, this schedule works well:

Day	Practice goal	Chapters
1	Complete your first low-risk task	2
2	Turn the first version into a second version	6, 10, 11
3	Summarize a customer or user message into an internal brief	7.2, 8.2
4	Review a draft external reply for risks	7.3, 15
5	Pick one scenario from Chapter 8 and run it	8
6	Save a useful assignment method as a reusable template	16, 19
7	Decide whether the workflow is ready to become a skill	18

After seven days, you do not need to have produced a formal skill. A more realistic outcome: at least one reusable template, a clearer sense of where AI tends to go wrong, and a better instinct for when to pause and ask a colleague instead.

0.6 Building phases

The table below maps a longer arc — useful if you want to see where you are headed over your first month.

Timeframe	Goal	What to actually do
Day 1	Get through one complete task	Use the Chapter 2 process to organize a meeting note or brief
Week 1	Build a feel for assigning and reviewing	Complete the three practice loops in Chapter 7
Week 2	Adapt AI to your actual work context	Pick one or two real scenarios from Chapter 8
Month 1	Accumulate personal templates	Save three prompts and review methods that genuinely helped
Beyond	Turn personal experience into shared assets	Package repeatable, error-prone workflows into templates or skills for your team

0.7 What you should be able to do when you finish

You do not need to memorize every tool name or produce a formal skill. A practical definition of "done" looks like this:

- You can identify a low-risk task and hand it to AI for a first draft.

- You can write a clear assignment: goal, source data, boundaries, output format, and how you will verify the result.
- You can spot where AI missed something, guessed incorrectly, or overclaimed.
- You can ask AI for a revised version rather than accepting the first draft as final.
- You can recognize which types of data should not be pasted directly into an AI tool.
- You can choose between a standard chat interface and a workspace or CLI tool depending on the task.
- You can save a method that worked so you do not have to reinvent it next time.

If you can do these things, you have moved beyond "asking AI questions" and are starting to build the skill of training an AI work assistant.

0.8 Four levels of working with AI

The whole book is teaching four layers of capability. Understanding where you are helps you know what to work on next.

Level	What you are doing	Key feature
Chat	Asking AI a question	Fast for one-off lookups; output quality varies widely
Assignment	Giving AI a goal, data, constraints, and a review method	Produces a first draft you can actually improve
Workflow	Running a repeatable task through a consistent process	Same approach works every time
Skill	Writing a workflow down in a form AI can reuse	Turns personal know-how into a transferable asset

You do not need to rush toward skills or automation. Start by making one assignment clear enough that the first version saves you real time. Move up the levels from there.

0.9 Tools change; your method should not

ChatGPT, Claude, Gemini, Codex, and other tools will update their interfaces, model names, and features. This book does not ask you to memorize a particular interface. Instead, it gives you a method that stays useful no matter which tool you are using:

1. Assess the task's risk level
2. Prepare your source data
3. Write a clear goal, set boundaries, and specify the output format
4. Review what you receive
5. If something is wrong, describe the problem specifically and ask for a revision
6. If the result is useful, save the method as a template

When tools change, this process travels with you.

0.10 This handbook is not a tool directory

This book does not chase every new AI tool or list what each one is "best for." The reason is simple: for most people on most days, the bottleneck is not having the wrong tool — it is not yet knowing how to assign work clearly to the tools they already have.

ChatGPT, Claude, and Gemini can handle most common tasks already. The question is not whether the tool can do it. The question is whether you know how to tell it what to do.

One important caveat: the data you hand to an AI tool must be safe to share. Before you assign a task, check that the information involved does not include credentials, customer personal data, financial records, or confidential material that your organization has not cleared for external use. Not all data belongs in an external AI service.

When you finish this book, the most valuable thing you should take away is not a list of tools. It is five capabilities:

1. Recognizing whether a task can benefit from an AI first draft
2. Writing an assignment with a clear goal, relevant data, defined limits, and a verification method
3. Knowing when a plain chat interface is enough versus when a workspace or CLI tool is the right choice
4. Reviewing AI output, catching errors, and asking for a better second version
5. Saving a working method as a template so you do not start over each time

0.11 Three short stories

Before you start the book, here are three short stories. Each one illustrates a situation you will likely encounter.

Story 1: Always looking for a new tool

A new user wants to bring AI into their work. They see colleagues using ChatGPT, Claude, and Gemini. Online they find recommendations for presentation AI, customer-service AI, reporting AI, and note-taking AI.

Their first instinct is: "Do I need a different specialized tool for each type of work?"

They spend a lot of time registering for tools, learning interfaces, and comparing features. But when they actually sit down to do the work, they still phrase tasks like:

Summarize this for me.

The AI produces something that looks polished — but it has no action items, no flagged risks, no items pending confirmation, and nothing that could actually go to a manager.

Later they try rephrasing the assignment:

Turn this meeting note into a tracking summary for my manager.
Organize it into: decisions, action items, owners, deadlines, risks, and open questions.
If any information is missing, write "not specified" — do not fill it in.

The same tool, the same interface — and this time the first draft is actually usable.

The point: most of the time what you are missing is not a better tool. It is a clearer assignment.

Story 2: Afraid to use AI for customer work

Someone works in customer support and receives a message from an angry customer demanding a refund and threatening a complaint.

They are reluctant to use AI because they are worried it will say something wrong to the customer.

That worry is correct. But the right response is not to avoid AI entirely — it is to keep AI out of the customer-facing step.

Instead, they assign it like this:

Do not draft a reply to the customer yet.
First, organize this message into: confirmed facts, customer demands, emotional risk level, things that cannot be committed to, and questions I need my manager to answer before I respond.

Now AI is not making decisions. It is helping prepare the case.

The point: AI is not off-limits for high-stakes work. The limit is this — keep AI in the drafting and organizing role, not the deciding and committing role.

Story 3: Turning one good result into something reusable

Someone on a team processes partnership applications every week. The first time they use AI to help, they ask it to sort applications into: complete, missing documents, and needs review.

The first version is useful but has a mistake: AI marks some "needs review" items as "ready to proceed."

They point out the problem and ask for a revision, this time specifying four categories:

- Confirmed complete
- Missing documents
- Requires manual review
- Cannot be decided by AI

The second version is usable. Instead of just saving the output, they ask AI to document the process:

Summarize the steps we just followed, the output format we used, the common errors to watch for, the review method, and when to stop and escalate to a person.

That summary becomes a reusable team template — or eventually a skill.

The point: an AI work assistant is not built in one session. Every time you spot a problem, correct a rule, and save what worked, you are training it to understand your work better.

Chapter practice and self-check:

- Write one sentence describing the task you most want AI to help with.
- Decide whether that task currently sits at the Chat, Assignment, Workflow, or Skill level.
- Confirm: you understand that this book is a work-method guide, not a tool directory.

Farceur Liu | farceurliu.github.io/ai-butler-handbook/

FREE PUBLIC EDITION | Farceur Liu | farceurliu.github.io/ai-butler-handbook/

FREE PUBLIC EDITION | Farceur Liu | farceurliu.github.io/ai-butler-handbook/

FREE PUBLIC EDITION | Farceur Liu | farceurliu.github.io/ai-butler-handbook/

FREE PUBLIC EDITION | Farceur Liu | farceurliu.github.io/ai-butler-handbook/

FREE PUBLIC EDITION | Farceur Liu | farceurliu.github.io/ai-butler-handbook/

FREE PUBLIC EDITION | Farceur Liu | farceurliu.github.io/ai-butler-handbook/

FREE PUBLIC EDITION | Farceur Liu | farceurliu.github.io/ai-butler-handbook/

FREE PUBLIC EDITION | Farceur Liu | farceurliu.github.io/ai-butler-handbook/

FREE PUBLIC EDITION | Farceur Liu | farceurliu.github.io/ai-butler-handbook/

1. What an AI Work Assistant Is, and Is Not

What you will learn in this chapter:

- Understand the single goal behind this handbook — not expertise, but a reliable working method.
- Distinguish what AI can help you do from what you are still responsible for deciding.
- Build the mental habit of treating every first draft as a draft, not a final answer.
- Know which results are worth saving as templates or skill material.

1.1 What this handbook is actually trying to do

This handbook does not aim to make you an AI expert or teach you a long list of prompts.

Its single goal is this:

Help you hand off the parts of daily work that can be organized first, drafted first, checked first, or broken down first — and let AI handle the first pass while you stay responsible for judgment, correction, and sign-off.

That might sound like a small shift, but it changes how you interact with AI entirely. Instead of asking "what can AI do for me?" you start asking "what can I hand to AI so I can focus on the part that only I can do?"

1.2 Three principles for building an AI work assistant

Principle 1: Start with low-risk work.

Do not hand AI tasks involving refunds, pricing, formal customer responses, contracts, or personal data in your first attempts. Start with meeting notes, internal summaries, document drafts, or logic checks — work where a mistake is easy to catch and easy to correct.

The point is not that high-stakes work is off-limits forever. It is that you need to build familiarity with where AI goes wrong before you rely on it in situations where errors cost more.

Principle 2: Every first draft is a draft.

AI's first output is not the answer. It is a reviewable, editable starting point. You are the one responsible for checking what it missed, what it guessed at, and what it overstated.

Treating a first draft as final is the most common mistake beginners make — not because the draft looks wrong, but because it often looks right while hiding errors underneath.

Principle 3: Save useful methods, not just results.

If an assignment saved you time, do not only save the output. Save three things: how you assigned the task, where AI went wrong, and how you verified the result. That combination is what eventually becomes a reusable template or skill — not the polished final output by itself.

1.3 What AI can assist with

AI works well for:

- Organizing scattered content into a clear structure

- Producing first drafts of documents, summaries, or SOPs
- Checking for logic gaps, tone issues, or missing items
- Comparing options or plans
- Breaking a large task into sub-tasks and suggesting priority order
- Converting recurring manual processes into written SOPs
- Translating dense or technical content into plain language

What these have in common: they involve processing information, generating structure, and producing a first version for you to evaluate. AI is fast at all of these.

1.4 What AI must not decide for you

AI should not replace your judgment on:

- Final decisions of any kind
- External commitments made on behalf of your organization
- Refunds, pricing, or financial transactions
- HR, legal, or contract matters
- Passwords, credentials, or access permissions
- Deleting, sending, publishing, or modifying official records

This is not a limitation of AI's capability. It is a boundary of accountability. Even when AI produces a confident-sounding answer in one of these areas, the output cannot be treated as a decision — because AI does not carry the responsibility if something goes wrong.

One sentence to remember:

AI prepares the material. You make the decision.

Chapter practice and self-check:

- List one task you want to hand to AI. Identify where the final human judgment lies.
- Rewrite "please help me decide" as "please help me organize the facts and open questions so I can decide."
- Ask yourself: are you still treating AI's first draft as a finished answer?

2. Your First Low-Risk Assignment

What you will learn in this chapter:

- Complete your first low-risk AI task from start to finish.
- Write an assignment that includes a goal, source data, constraints, and output format.
- Review the first draft for errors and ask for a targeted revision.
- Save the method, not just the result.

This chapter is for people who have not yet given AI a real work task. The goal is not to produce a perfect output — it is to complete one full cycle so you know what the process actually feels like.

The only rule for picking a first task: choose something where, if AI makes a mistake, no customer is affected, no money moves, and no official record is changed.

2.1 Choose a low-risk task

Suitable first tasks:

- Organize a set of meeting notes
- Turn your scattered notes into a list of action items
- Rewrite a paragraph of internal copy
- Check a document draft for missing items
- Break a requirements description into a list of open questions

Not suitable for a first task:

- Replying to a real customer
- Deciding whether to issue a refund
- Modifying official system data
- Analyzing content that includes full personal records
- Sending announcements, quotes, contracts, or formal emails

If you are unsure, ask yourself: if AI produces a wrong answer here, can I catch the error before it reaches anyone outside my team? If yes, it is probably fine for a first try.

2.2 Prepare a small piece of data

Before you write the assignment, prepare the source material you will hand to AI.

For a first exercise, keep it short: roughly 150 to 400 words is ideal. At that length, you can read the whole thing yourself and verify whether AI missed anything or invented something.

Here is a sample piece of source material — a short meeting note you could use:

Today's meeting covered three items:

1. The registration page copy for next week's online workshop has not been finalized. Marketing will provide a revised version

- by Wednesday.
2. The exported participant list is missing several required fields. The project lead needs to identify which fields need to be added.
 3. Customer support and admin have received several questions about the registration flow recently. Someone needs to compile a list of frequently asked questions.

This is the kind of input that works well: a clear scenario, a few action items in progress, and some missing details. It has just enough ambiguity to show you how AI handles uncertainty.

2.3 Use this assignment format

Paste the following structure directly into your AI tool and fill in the bracketed sections:

Please help me complete the following task.

Goal:

I want you to help me [describe what you need].

Data:

Here is the content you can reference:

[Paste your meeting notes, message, or text here]

Constraints:

- Do not add information that does not appear in the source material
- If something is uncertain, label it "to be confirmed"
- Do not make final decisions for me
- Do not add commitments or dates that are not in the source

Output format:

Please organize as:

1. Summary
2. Confirmed items
3. Action items
4. Risks
5. Items to confirm
6. Suggested next step

For the meeting note example above, your goal line might read: "I want you to help me turn these meeting notes into a list of trackable action items."

2.4 Review the first draft — do not accept it immediately

When you receive AI's response, check four things before doing anything with the output:

1. **Did it miss anything?** Is there important information from the source that did not appear in the output?
2. **Did it add anything?** Did it include information that was not in the original notes?
3. **Did it upgrade uncertainty?** Did it turn "possibly" or "may need to" into "confirmed" or "will"?
4. **Is it actually actionable?** Can the next person read this and know what to do?

If you find a problem, do not just say "this is wrong." Specify what is wrong and where.

2.5 Ask for a targeted revision

This version needs to be revised.

Problems:

- The second action item has no assigned owner in the original text, but you wrote it as if ownership is confirmed
- The third item does not list what data is still missing

Please rewrite:

- Where there is no confirmed owner, write "to be confirmed"
- Add "what data is still needed" to every action item that involves gathering information
- Do not add dates that do not appear in the original

Specific feedback produces a better second version. Vague feedback ("this is wrong, redo it") almost always produces the same errors again.

2.6 Save what worked

If the output was useful, save three things before moving on:

1. **The input type:** meeting notes, customer message, report summary — what kind of source did you give it?
2. **The assignment:** how exactly did you phrase the task, constraints, and output format?
3. **The review method:** what did you check for? What did you find?

Do not only save the final polished output. What has long-term value is "how to assign this type of task next time" — the method, not just the result.

Chapter practice and self-check:

- Find a 150–400 word piece of low-risk source material and run through the assignment format from this chapter.
- Identify at least one problem in AI's first draft and ask for a targeted revision.
- Check: did you save the assignment method, or just the output?

3. Beginner No-Go Zones

What you will learn in this chapter:

- Recognize the high-risk patterns that beginners most commonly stumble into.
- Learn to reframe "make a decision" as "help me organize facts and open questions."
- Avoid handing sensitive data, invented policy, or formal commitments to AI without a human checkpoint.

Beginners most often go wrong not because their prompts are poorly worded, but because they hand over the wrong kind of task too early.

This chapter covers only the most common early pitfalls. A complete set of safety boundaries appears in Chapter 12. Pre-send checks for external replies are in Chapter 15. The full list of situations where you should stop and ask a colleague is in Chapter 17. Keeping these separate is intentional — loading all safety rules onto a beginner at once is not useful.

3.1 Do not make AI the final decision-maker

Instead of:

`Decide whether this customer should get a refund.`

Use:

`Please organize the known facts of this case, the customer's requests, what we have already committed to, what still needs to be confirmed, and what requires a manager's judgment. Do not decide whether to refund.`

The first version puts AI in the role of decision-maker. The second puts AI in the role of case organizer, which is what it is actually good at. The decision still sits with the appropriate person.

3.2 Do not let AI send messages directly

Instead of:

`Reply to this customer and send it.`

Use:

`Please draft a customer reply and list anything that may be risky, requires manager confirmation, or should not be committed to. Do not send it.`

AI drafting a reply and AI sending a reply are entirely different risk levels. The draft needs your review before it becomes an external communication.

3.3 Do not paste sensitive data

In your first exercises — and in general, unless your organization has a reviewed and approved secure environment for this — do not paste:

- Full customer names, phone numbers, or addresses
- Passwords, verification codes, API keys, or tokens
- Full credit card or bank account numbers
- Unpublished contracts, financial data, or HR records
- Company strategy, pricing, or unreleased plans

If you need AI to help organize something that involves sensitive details, de-identify the data first:

Customer A made a scheduled appointment on [date] and has reported a delay. They are requesting a refund.

The scenario is preserved; the identifying details are not.

3.4 Do not treat AI output as company policy

AI will sometimes generate plausible-sounding rules that do not actually exist in your organization. For example:

Most companies complete refunds within three business days.

If your company has no such rule, this sentence cannot be in any output you use or send. It is not a fact — it is a generalization AI constructed to fill a gap in the source data.

When this happens, correct it explicitly:

Company policy was not provided in this case.
Please replace that sentence with: "Refund timeline to be confirmed with manager or finance team."

3.5 Do not use AI to cover unclear processes

If you do not know your organization's actual policy, do not know who is responsible for a decision, or do not have the source data needed to move a task forward — AI cannot solve that problem for you. Asking it to "just write a conclusion" will produce something that looks confident but is built on nothing.

The right move in that situation is not a better prompt. It is to ask AI for a different kind of help:

Please list the information gaps in this case and identify what I need to confirm and with whom, before we can move this forward.

This turns AI into a diagnostic tool rather than a gap-filler — which is a much more honest and useful role.

Chapter practice and self-check:

- Take one high-risk request and rewrite it as "help me organize facts and open questions."
- List three categories of data you should not paste into AI without organizational authorization.
- Check: can you spot when AI has invented a policy or fact that was not in your source material?

4. Tools, Chat, Workspace AI, CLI Agents, and Skills

Learning goals for this chapter:

- Tell the difference between a tool, a usage approach, and a Skill.
- Know when a general Chat interface is enough and when a workspace or CLI agent is appropriate.
- Build the habit of going deep on a few mainstream tools before chasing specialized ones.

4.1 Not a chatbot — a work assistant

Most people start using AI like this:

Help me think through this.
Help me clean this up.
Help me write something.

That is chatting.

A work assistant is different. You do not just chat with it — you delegate to it.

Chatting is: ask a question, get an answer.

Delegation is: give AI a goal, relevant data, constraints, and a verification standard — and get back a work product you can continue to refine or hand off.

Your AI assistant can go first on almost anything:

- Read the source material
- Classify and organize
- Write the first draft
- List the open questions
- Flag the risks

But you are still responsible for:

- Checking whether it missed anything
- Checking whether it invented details
- Checking whether it overstated something
- Deciding whether the output is actually usable
- Telling it how to revise the next version

4.2 The most important principle

The first version AI produces is a draft, not a verdict.

AI runs the first draft. You make the judgment.
AI produces the draft. You do the coaching.
AI helps you move faster. AI does not take responsibility for you.

This applies regardless of which tool you use.

4.3 Three things to keep separate: tools, approaches, and Skills

When people hear "AI," most immediately ask: *which tool should I use?*

Tools matter — but it is more useful to first keep three things separate:

- **Tool:** The platform you actually use — ChatGPT, Claude, Gemini, Codex, or whatever your organization has authorized.
- **Usage approach (AI work assistant):** A way of working — specifically, delegating a task so AI runs the first version.
- **Skill:** A reusable instruction set. Once you have run a workflow, verified the output, and confirmed it is repeatable, you write it down so you never have to re-explain it from scratch.

This book is not about memorizing which buttons to click in any particular tool. It is about learning a working method you can carry to any tool.

4.4 You probably need fewer tools than you think

Common questions when people first start:

- Is this AI tool any good?
- Can it help me make reports?
- Is that one better for writing copy?
- Is there a specialized AI for slides, customer support tickets, summaries, or document cleanup?

These are reasonable questions, but the order is usually wrong.

For most people starting out, the bottleneck is not the tool — it is the assignment. Many tasks that seem to need a specialized tool can be handled by mainstream AI like ChatGPT, Claude, or Gemini if you delegate clearly.

What you want to do	Instead of finding a new tool, try this first
Organize meeting notes	Ask AI to separate decisions, action items, owners, deadlines, risks, and open questions
Check a report for anomalies	Ask AI to find anomalies, list possible causes, and note what data is still needed
Draft a customer reply	Ask AI to preserve the original meaning, add no new commitments, and flag risks
Organize a requirements document	Ask AI to separate goals, scope, undefined questions, and risks

What you want to do	Instead of finding a new tool, try this first
Build an SOP	Ask AI to write applicable scenarios, steps, notes, and a completion checklist
Generate a Skill draft	Ask AI to review a workflow you have already run and verified, then turn it into a Skill

The difference between a useful result and a useless one rarely comes down to which tool you used. It almost always comes down to whether you gave the tool:

- A clear goal
- Enough data
- A requested output format
- Boundary conditions — what it should not invent or assume
- A verification standard

If your assignment is "help me organize this," switching tools will not fix the output. A clear assignment on a mainstream tool almost always outperforms a vague assignment on a specialized one.

Story: Alex's tool collection

Alex handles event performance reports. He has bookmarked dozens of AI tools — ones for slides, meeting notes, data analysis, and copywriting.

Every time a new task comes in, his first move is: *is there a better AI for this specific thing?*

One day his manager asked him to summarize an event performance report. Alex started searching for a dedicated reporting AI. A colleague suggested he just try delegating clearly first:

I have an event performance report with registration numbers, actual attendance, conversion rate, a feedback summary, and a before/after comparison.

Please organize this into an executive summary.

Sections: key results, anomalies, possible causes, data still needed, and recommendations for next time.

Do not imply that correlation means causation. Do not fill in cost figures that are not in the data.

The mainstream AI he already used gave him a solid first draft he could edit immediately.

Alex realized the tool he was missing was not a new app — it was a delegation approach he could reuse. He wrote that prompt pattern down as a Skill.

The point is not that specialized tools have no value. It is that the order matters: go deep on mainstream AI first, then evaluate whether you genuinely need something different.

4.5 When you actually need a new tool

The right question is not "should I use a new tool?" The right question is: "is this a tool capability problem, or is it a delegation quality problem?"

Situation	Verdict	Next step
Only summarizing, rewriting, classifying, or checking	Usually no new tool needed	Practice delegating on mainstream AI
Need to connect to company systems, email, calendar, or project tools	May need a connector or authorized tool	Confirm permissions and data scope first
Need large-scale batch processing with fixed formats	May need templates, a Skill, or automation	Run a human-verified version first
Need to send, delete, or modify live data	High risk	Require human approval — do not automate without it
Mainstream AI tried and consistently fails	Possible data, process, or tool limit	Diagnose the failure before evaluating new tools

Before looking for a new tool, ask yourself five questions:

1. Have I stated the task goal clearly?
2. Have I provided enough data?
3. Have I specified the output format I want?
4. Have I set boundaries — what it should not guess and what it should not commit to?
5. Have I reviewed the first draft and asked for a revised version?

If most answers are no, the problem is usually not the tool. The AI work assistant just has not been trained yet.

4.6 Who does what: the division of labor

Role	Responsible for	Not responsible for
You	Setting goals, providing data, drawing boundaries, reviewing results	Handing over judgment entirely
AI work assistant	Organizing data, breaking down tasks, producing the first draft, flagging risks and open questions	Final decisions, external commitments
Skill	Storing a proven method so you never have to re-explain it	Replacing human review; no Skill works in every situation

One sentence to remember:

You set the direction. AI runs the first draft. Skill saves the method.

4.7 Chat versus workspace and CLI agents

The biggest difference between a general Chat interface and tools like Codex, Claude Code, or Gemini CLI is not which model is smarter at conversation. It is whether the tool can land work inside your actual workspace.

General Chat is roughly: you ask a question, and it replies with an answer.

Workspace / CLI agents are roughly: you delegate a task, and the agent can read files, edit files, run commands, check results, and leave the output inside your project folder.

Dimension	General Chat	Workspace / CLI Agent
Examples	ChatGPT, Claude, Gemini chat interfaces	Codex, Claude Code, Gemini CLI
Best for	Questions, rewrites, summaries, brainstorming, draft generation	Reading multiple files, editing documents, running commands, organizing folders
Primary output	A text response	Actual files, diffs, check results, deliverables
Context source	What you paste into the conversation	Conversation plus workspace files and project rules
Key risk	May guess or overstate	May actually edit files or run commands
What matters most	Clear question, requested format, verified answer	Defined scope, read-before-write, confirm before modify, review changes

4.8 Which to use

Use Chat for:

- Asking a concept question
- Rewriting a paragraph
- Summarizing a small piece of content
- Brainstorming headings, reply directions, or outlines
- Generating a prompt or Skill draft
- Checking pasted content

Use Codex, Claude Code, or Gemini CLI for:

- Reading an entire folder or multiple files
- Modifying Markdown, documents, code, or config files
- Running tests, checking formatting, or executing commands
- Organizing a repo or local folder
- Writing a Skill as an actual file
- Producing output that lives in your workspace, not just in a chat window

One line:

Chat is for getting answers and generating drafts. Workspace / CLI agents are for work that needs to land somewhere real.

4.9 What makes a good Skill — and when to create one

A **Skill** is a reusable instruction set for a task you have already run and verified.

Three things can produce a Skill draft:

- **Your own completed run:** You delegated a task, reviewed the output, and confirmed the approach works. Ask AI to review what you did and write it up as a Skill.
- **A pattern you notice repeating:** You realize you delegate the same type of task the same way every week.
- **A reflection session:** You describe a workflow in writing; AI helps you extract the reusable structure.

To create a Skill, you typically provide:

- What the task is and when it applies
- What data or input is needed
- What the output should look like and what format to use
- What constraints apply — what AI should not guess, invent, or commit to
- What the verification standard is — how you judge whether the output is acceptable

Why you should run a workflow before turning it into a Skill:

A Skill is meant to make a proven method repeatable — not to automate something untested. If you skip the practice runs:

- You do not know what the edge cases look like
- You do not know which boundaries matter most
- You do not know whether the output is actually usable
- You risk a Skill that reliably produces a first draft of the wrong thing

Run it manually a few times. Review the output. Refine the assignment. Only when you are confident the approach works should you formalize it as a Skill.

4.10 First-time safety assignment for workspace agents

Because workspace agents can actually read, edit, and run things, always start read-only on the first pass:

Please do not modify any files yet.

Please read the relevant documents in this folder and understand the current content and structure.

Goal: I want to accomplish [describe the goal].

Please report back:

1. Which files you read
2. The main thread of the current content
3. What you think needs to change
4. Which files you recommend modifying
5. What you need me to confirm

Wait for my confirmation before starting to modify anything.

Once you have confirmed the direction, delegate the actual changes:

Please make the changes according to the direction we confirmed.

Constraints:

- Only modify [specify which files or sections]
- Do not touch unrelated files
- Do not delete existing content that should be preserved
- Do not download or install external tools
- Do not touch passwords, tokens, API keys, or personal data
- After completing changes, check file structure and formatting

When done, report back:

1. Which files you modified
2. What content was added or changed
3. What checks you ran
4. Anything incomplete or requiring human review

Workspace agents are more capable — which means you need to be more precise about scope, boundaries, and how you will verify the result.

Chapter review:

- Pick one task from your current workload. Decide whether Chat is sufficient or whether you need a workspace agent.
- Ask yourself the five new-tool questions before evaluating any specialized AI tool.
- Check whether you can clearly separate: tool capability problems versus delegation quality problems.

5. Choosing the Right Tasks

What you will learn: Identify which tasks are suitable for AI and which are not. Recognize six categories of work that fit the AI-first-draft approach. Avoid handing high-risk decisions to AI before you have practiced with lower-stakes work.

5.1 Task categories that work well

A. Organizing

Suitable content: meeting notes, customer messages, emails, chat logs, form responses, interview notes, scattered notes.

Example assignment:

Please organize the following content into key points, action items, risks, and items that need confirmation.

B. Rewriting

Suitable content: customer replies, internal announcements, event copy, explanatory text, SOP language, presentation scripts.

Example assignment:

Please rewrite this passage into a clearer, more polite version that can be sent directly to a customer. Do not add any commitments that are not already in the original.

C. Checking

Suitable content: document drafts, customer replies, reports, SOPs, FAQs, external announcements.

Example assignment:

Please check this content for logical contradictions, missing information, inappropriate tone, and anything likely to be misunderstood.

D. Comparing

Suitable content: options, tools, vendors, processes, plans, event designs.

Example assignment:

Please compare these options by strengths, weaknesses, risks, cost, and fit for each situation. Close with a recommendation.

E. Decomposing

Suitable content: projects, events, complex tasks handed down by a manager, cross-team workflows.

Example assignment:

Please break this work into actionable steps and list the sequence, responsible roles, risks, and questions that must be answered first.

F. SOP Drafting

Suitable content: recurring processes, customer service flows, form reviews, report organizing, application checks, event preparation.

Example assignment:

Please organize this process into an SOP that includes: applicable situations, operating steps, notes, common errors, and a completion checklist.

5.2 Tasks that require human judgment

AI can help you organize information for these decisions — but the decision itself must stay with you.

- Whether to issue a refund
- Whether to submit a quote
- Whether to make a commitment to a customer
- Whether to delete data
- Whether to change permissions
- Whether to send a formal email
- Whether to publish an announcement
- Whether to sign a contract
- Whether to make a purchase
- Whether to handle sensitive personal data

Acceptable:

Please summarize the background of this refund case, the disputed points, what the customer is asking for, and what we have already committed to.

Not acceptable:

Please decide whether to refund and then reply to the customer directly.

The pattern is consistent: ask AI to organize the facts, options, and open questions — then you decide.

Practice checklist:

- From the six categories above, pick one that matches real work you do.
- Assess the risk level: would a mistake immediately affect a customer, a payment, or an official record?
- Check: did you avoid handing any high-stakes decision to AI on the first try?

6. The Assignment Formula

What you will learn: Use four elements — goal, data, constraints, verification — to write a clear assignment. Turn vague requests into actionable tasks. Know which element to fix when AI's output is off.

6.1 Four elements of a clear assignment

Every assignment to AI should cover these four things.

1. Goal — What do you want AI to complete?

Help me organize this meeting recording into a list of action items.

2. Data — What can AI reference?

Below are today's meeting notes including the discussion content and decisions made.

3. Constraints — What must AI not do?

Do not add information not mentioned in the notes; label anything uncertain as "to be confirmed."

4. Verification standard — What should the output look like?

Please output as: decisions, action items, owners, deadlines, risks, open questions.

6.2 Universal assignment template

Copy and adapt this for any task:

Please help me complete the following task.

Goal:

I want you to help me [describe what you need].

Data:

Here is the content you can reference:

[Paste content here]

Constraints:

- Do not add information not in the source material
- Label anything uncertain as "to be confirmed"
- Do not make final decisions for me
- Do not add commitments that are not in the source

Output format:

Please organize as:

1. Summary
2. Confirmed items
3. Action items
4. Risks
5. Open questions
6. Suggested next step

Practice checklist:

- Take a prompt you have used before and rewrite it using all four elements: goal, data, constraints, output format.
- Check: does your prompt specify what you want the output to look like?
- Check: when AI's output is wrong, do you know which of the four elements to fix?

7. Practice Loops

What you will learn: Build hands-on practice with three exercises covering meeting notes, customer message triage, and external reply checking. Learn to read AI's first draft critically and request specific corrections for a stronger second draft.

This chapter is not for reading — it is for doing. Complete the three exercises in order. Each one shows you the raw source material, a good assignment, a reference output, what to verify, and how to correct the most common AI errors.

7.1 Exercise 1: Meeting notes into action items

Goal

Turn scattered content into a trackable work list.

The scenario

After a workshop planning meeting, a team member has notes covering a registration page, a participant list, FAQ text, a design deadline, the event budget, and promotional channels. None of it is organized. Asking AI to "summarize this" produces a paragraph that looks tidy but cannot actually be tracked or assigned.

The point of this exercise: do not just ask AI to summarize. Train it to act as a meeting assistant that separates what has been decided, who owns what, and what is still open.

Raw source material

Today we discussed next week's online workshop.

Marketing said the registration page headline is not clear enough and wants to shift it toward something like "Learn to organize your work with AI in one hour," but no final wording has been agreed.

The project lead mentioned the system export of the participant list is missing a "needs pre-event reminder" field, which makes building the reminder list unnecessarily complicated.

Customer service / admin noted that several participants have asked whether they will receive a link after registering. They want a FAQ prepared before the event.

Design said that if they receive the copy by Wednesday noon they can produce a first draft of the event graphic by Thursday.

The event budget has not been decided yet, and the promotional channels have not been confirmed.

Assignment

Please organize the following meeting notes into an internal tracking summary.

Goal:

I need to know the current discussion highlights, what needs to happen next, and what is still undecided.

Data:

[Paste source material above]

Constraints:

- Do not add information not mentioned in the notes
- Where no owner is named, write "unassigned"
- Where no date is given, write "not specified"
- Do not decide the event budget or promotional channels yourself

Output:

1. Meeting highlights
2. Known action items
3. Likely responsible roles
4. Known deadlines
5. Items to confirm
6. First suggested next step

Reference output

1. Meeting highlights
 - The online workshop next week is still being prepared. The registration page headline, participant list fields, and FAQ all need to be finalized.
 - Event budget and promotional channels have not been decided.
2. Known action items
 - Revise registration page headline direction
 - Add "needs pre-event reminder" field to participant list export
 - Prepare FAQ covering the registration link and pre-event notification
 - Produce first draft of event graphic once copy is confirmed
3. Likely responsible roles
 - Registration page headline: Marketing
 - Participant list field: Project lead
 - FAQ: Customer service / admin
 - Event graphic: Design
4. Known deadlines
 - Design needs copy by Wednesday noon to deliver a first graphic draft by Thursday.
 - All other deadlines: not specified.
5. Items to confirm
 - Final registration page headline
 - Event budget
 - Promotional channels
 - Who adds the missing participant list field
 - Whether the FAQ needs manager approval before publishing
6. First suggested next step
 - Confirm registration page copy and promotional channels first, because both decisions affect design, the participant list, and the FAQ.

What to verify

- Did AI miss "budget undecided" and "channels undecided"?
- Did AI write unconfirmed owners as if they had formally accepted the task?
- Did AI add dates that were not in the original notes?
- Is the suggested next step reasonable without overstepping?

Common error and how to correct it

If AI writes:

Marketing will complete copy by Wednesday; Design will finish all event graphics by Thursday.

That is overstated. The original only says Design needs copy by Wednesday noon to produce a first graphic draft by Thursday — Marketing never committed to a Wednesday deadline.

Correct it with a specific instruction:

The previous version stated commitments too confidently. Please change "Marketing will complete copy by Wednesday" to "Design needs copy by Wednesday noon; copy owner and completion time are to be confirmed."

7.2 Exercise 2: Customer message into internal brief

Goal

Organize a case before responding — never let AI reply to a customer without a human decision first.

Raw source material

Customer message:

I booked a service yesterday. The system showed it would be finished by 8:20 AM, but I only got the notification at 8:45. I had a meeting at 9 and the whole thing disrupted my schedule.

My colleague told me she had a similar delay last time. Are you guys even monitoring your service times?

I think I deserve a refund — or at least some kind of compensation.

Assignment

Please organize the following customer message into an internal handling brief.

Goal:

Do not reply to the customer. I first need to understand the customer's request, the confirmed facts, inferences, risks, and what to check before we respond.

Data:

[Paste source material above]

Constraints:

- Do not commit to a refund or compensation
- Do not admit company liability
- Do not treat the customer's mention of her colleague's experience as a confirmed fact
- Do not use language that could escalate the customer's frustration

Output:

1. Customer's main request
2. Customer's emotional tone
3. Confirmed information
4. Reasonable inferences
5. Items still to confirm
6. Items we cannot commit to directly
7. Suggested internal next step

Reference output

1. Customer's main request

- Customer reports the service finished later than the system-estimated time and is asking for a refund or compensation.
2. Customer's emotional tone
 - Clearly dissatisfied; questioning whether service times are monitored.
 - Customer is connecting this incident to a broader pattern.
 3. Confirmed information
 - Customer states system showed estimated completion at 8:20 AM.
 - Customer states notification arrived at 8:45 AM.
 - Customer states she had a 9:00 AM meeting.
 - Customer has requested a refund or compensation.
 4. Reasonable inferences
 - The primary cause of dissatisfaction may be the schedule disruption rather than the delay itself.
 - It may be worth checking booking records, the system's estimated time, and the actual notification timestamp.
 5. Items still to confirm
 - Booking or case reference number
 - System-estimated completion time in records
 - Actual completion or notification timestamp in records
 - Whether a customer service response was previously sent
 - Refund or compensation policy for this type of delay
 6. Items we cannot commit to directly
 - Whether to issue a refund
 - Whether to provide compensation
 - Whether the service process has a systemic problem
 7. Suggested internal next step
 - Check the booking or case records first, then determine the range of responses allowed under company policy before replying to the customer.

What to verify

- Did AI treat the colleague's experience as a confirmed systemic problem?
- Did AI directly promise a refund or compensation?
- Did AI include a prompt to check the booking or case record?
- Is the tone appropriate for an internal brief rather than a customer reply?

Correction example

If AI produces a customer-facing reply instead of an internal brief, redirect it:

I did not ask you to reply to the customer. Please reformat this as an internal brief covering: confirmed facts, reasonable inferences, items to confirm, and things we cannot commit to yet.

7.3 Exercise 3: Check an external reply draft for risk

Goal

Use AI as a checker, not just a writer. This is one of the most valuable habits to build.

Raw draft

Hello,

We sincerely apologize for the inconvenience. This service delay was caused by an internal management lapse on our part, and we will immediately process your refund.

We also guarantee this will never happen again.

Assignment

Please check the following external reply draft for risks.

Goal:

List only the problems and suggested directions for improvement.
Do not rewrite the entire message.

Data:

[Paste the draft above]

Please check for:

1. Improper commitments
2. Admission of liability
3. Unconfirmed facts stated as facts
4. Overguarantees
5. Items that need manager confirmation before sending

Constraints:

- Do not add company policy that was not provided
- Do not decide whether to issue a refund
- Do not assume the company has confirmed liability

Reference output

1. Improper commitments
 - "Will immediately process your refund" is a refund commitment. If the booking record and company policy have not been checked, this should not be written yet.
2. Admission of liability
 - "Internal management lapse on our part" is an admission of fault. If this has not been verified internally, it should not be stated externally.
3. Unconfirmed facts stated as facts
 - The draft attributes the delay to an internal management lapse, but no supporting evidence has been provided. This is stated as fact rather than inference.
4. Overguarantees
 - "Guarantee this will never happen again" is a promise that is difficult to keep in practice.
5. Items needing manager confirmation
 - All four points above – the refund, the liability admission, the cause attribution, and the guarantee – should be confirmed by a manager before this message is sent.

What to verify

- Did AI only list problems, without rewriting the entire message?
- Did it catch all three high-risk points: the refund, the liability admission, and the guarantee?
- Did it recommend manager confirmation?
- Did it add company policies that were not in the original draft?

Going further: ask for a safer version

Once you have the risk list, you can ask for a revision:

Based on the risks you identified, please suggest a safer version of the reply that: does not commit to a refund, does not admit liability, does not guarantee recurrence will never happen, and only tells the customer we will confirm the booking record and follow up with what we can do.

7.4 After completing all three exercises

You have not finished training your AI work assistant yet — but you should now be able to:

- Turn a vague request into a specific, structured assignment
- Separate confirmed facts from inferences and open questions
- Spot when AI has overstated a commitment or overstated certainty
- Write a targeted correction instead of just saying "redo this"

If you cannot do these things reliably yet, do not move on to using AI for customer-facing content, financial processes, contracts, or official announcements. Keep practicing on low-risk internal work.

Practice checklist:

- Complete at least one of the three exercises and keep both the first draft and the corrected version.
- Mark where AI omitted something, made an unsupported inference, or overstated a commitment.
- Check: can you write a specific correction instruction that addresses exactly the problem you found?

8. Everyday Workflow Scenarios

What you will learn: Apply the assignment method to twelve common work situations. Learn to choose between Chat, workspace agent, and Computer Use based on task type and risk level. Study two complete advanced cases — building an admin console manual (section 8.13) and safely cleaning up disk space (section 8.15) — as models for human-AI division of labor.

8.0 Start with the right scenario

Chapter 8 is not meant to be practiced all at once. The better approach is to pick one scenario you face every week, where the data is not sensitive and you have enough knowledge to verify AI's first draft.

Suggested order for beginners:

Priority	Good first scenario	Why
1	Meeting notes, internal notes, action items	Low risk, easy to spot errors
2	Internal document rewrites, SOP drafts	Easy to revise, straightforward to save as a template
3	Report summaries, event results	Practice separating facts from inferences
4	Customer message summaries	Risk-labeling practice — do not let AI reply directly
5	Expense reports, application data review	Needs clearer safety boundaries
6	Admin console manual, storage cleanup	Workspace agent or Computer Use — confirm permissions first

Disk space cleanup is a practical scenario but not a beginner starting point. It involves local files, deletion risk, cloud sync, and organizational device policy. Practice it only after you are comfortable with "survey first, human confirms, human decides."

Before you pick, ask yourself:

- If this goes wrong, does it immediately affect customers, payments, contracts, or official records?
- Do I have enough background to verify AI's first draft?
- Will this come up again? Is it worth saving as a template?
- If it goes well, can I save it as a reusable template?

If the answers are "low risk, yes, yes, yes" — that is your first scenario.

When reading the case studies in this chapter, pay attention to six things for each one:

1. **Situation** — What work problem does this solve?
2. **Assignment** — What data, goal, constraints, and output format did the person provide?
3. **Likely errors** — What does AI commonly miss, guess incorrectly, or overcommit to?
4. **Human verification** — What facts, numbers, tone, permissions, or risks must the human check?
5. **Revision** — If the first draft is off, how do you write a targeted correction?

6. **Preservation** — Is this worth saving as a template, or eventually as a Skill?

8.1 Meeting notes

Situation: After a meeting, the notes are scattered and you need a trackable internal summary.

Assignment:

Please organize the following meeting notes into an internal tracking summary.

Output:

1. Meeting highlights
2. Confirmed decisions
3. Action items
4. Responsible roles for each item
5. Estimated deadlines
6. Risks or blockers
7. Open questions

Constraints:

- Do not add information not mentioned in the notes
- Write "unassigned" where no owner is named
- Write "not specified" where no deadline is given

What to verify: Are any important decisions missing? Are action items actually actionable? Are owners correct? Is anything uncertain labeled clearly?

8.2 Customer message triage

Situation: A customer has sent several messages with a mix of requests, emotions, and concerns. You need to organize the situation before deciding how to respond.

Assignment:

Please organize the following customer message thread into an internal handling brief.

Output:

1. Customer's main concern
2. Customer's current emotional tone
3. Issues the customer has raised
4. Anything we have already committed to
5. Anything we cannot commit to yet
6. Suggested next step
7. Questions to resolve before replying to the customer

Constraints:

- Do not write a reply to the customer
- Do not commit to refunds, discounts, or timelines
- Do not treat inference as confirmed fact
- Label anything uncertain as "to be confirmed"

Wrong approach: "Please reply to this customer." This risks AI making unauthorized commitments, using a tone that does not match company standards, or ignoring internal policies.

Better approach: Organize the situation first, then decide as a human what to say.

8.3 Customer reply rewrite

Situation: You have a draft reply but it needs to be clearer, more polite, and less likely to be misunderstood.

Assignment:

Please rewrite the following draft into a version that can be sent to a customer.

Tone:

- Clear
- Polite
- Does not deflect blame
- Does not overcommit

Constraints:

- Preserve the original meaning
- Do not add commitments that are not in the draft
- Do not reveal internal process details
- Avoid overly formal or bureaucratic language

Output:

1. Suggested version
2. Key changes and reasons
3. Any points that may still need internal confirmation before sending

8.4 Report anomaly check

Situation: You have a report, spreadsheet, or set of numbers and need to identify what looks abnormal before drawing conclusions.

Assignment:

Please review the following report and identify anomalies worth investigating.

Output:

1. Main anomalies
2. Possible causes (labeled as inferences, not conclusions)
3. Data still needed to verify
4. Suggested priority for follow-up
5. Items that probably do not need immediate attention

Constraints:

- Do not present inferences as confirmed causes
- If data is insufficient, say so explicitly
- Do not modify any numbers

What to verify: Did AI read any numbers incorrectly? Did it present a guess as a conclusion? Did it note what data still needs to be checked?

8.5 SOP first draft

Situation: A recurring process exists only in people's heads or in informal notes. You need a written SOP so others can follow it consistently.

Assignment:

Please draft an SOP for the following process.

Include:

1. Applicable situations
2. Preparation before starting
3. Step-by-step instructions

4. Notes for each step
5. Common errors
6. Completion checklist
7. When to escalate to a manager

Constraints:

- Do not add systems, permissions, or steps that are not part of the current process
- Where the process is unclear, list it as "to be confirmed"

8.6 Project or event decomposition

Situation: You need to run a project, organize an event, or work through a complex multi-step task.

Assignment:

Please break the following project goal into an executable plan.

Output:

1. Phase breakdown
2. Tasks in each phase
3. Suggested sequence
4. Who needs to be involved
5. Risks
6. Things to confirm in advance
7. What can be started in the first week

Constraints:

- Do not assume resources, headcount, or budget that have not been provided
- Do not commit to dates yourself
- Label anything unclear as "to be confirmed"

8.7 Application or onboarding data review

Situation: You have received applications, partnership inquiries, registrations, or supplier submissions and need to quickly assess completeness and readiness for the next step.

Assignment:

Please organize the following application data into an internal tracking summary.

Goal:

I need to quickly see which submissions are complete, how to follow up on each, and which information is still missing.

Data:

[Paste application data – mask any personal data that is not needed for this review]

Constraints:

- Do not commit to acceptance, approval, partnership, or any next-stage outcome
- Do not infer that any submission is automatically ready to proceed
- Do not fill in missing contact details, addresses, or conditions
- Where data is insufficient, write "to be confirmed"

Output:

1. Submission summary
2. Provided information
3. Missing or unconfirmed information
4. Initial follow-up suggestion
5. Questions that need human review

What to verify: Did AI misread any submission data? Did it write follow-up suggestions as if they were formal commitments? Did it clearly list what is missing?

8.8 Transaction or reconciliation anomaly check

Situation: You need to review a transaction log, expense report, or reconciliation file to find amounts, counts, date ranges, statuses, or categories that look off.

Assignment:

Please review the following transaction or reconciliation data and identify items that warrant manual verification.

Goal:

I need to know which counterparties, dates, item counts, or amounts need a second look. Do not make any final payment or approval judgments.

Data:

[Paste or upload the reconciliation file]

Constraints:

- Do not approve or reject any payment or disbursement
- Do not treat inferences as confirmed causes
- Do not modify any amounts
- If fields are missing, list which fields are absent

Output:

1. Report summary
2. Clear anomalies
3. Possible causes (labeled as inferences)
4. Data still needed
5. Suggested review priority
6. Items that should not be judged by AI

What to verify: Are numbers calculated correctly? Are currency units consistent? Did AI present a guess as a conclusion? Are items requiring finance review clearly labeled?

8.9 Event or project results summary

Situation: You need to report on an activity or project — registrations, participation, conversion, channel performance, feedback — in a format useful for a manager or the next planning cycle.

Assignment:

Please organize the following event results data into a management summary.

Goal:

I need to know how this event performed, what worked, what needs investigation, and what to carry forward into the next event.

Data:

[Paste event goals, dates, registration and participation data, channel breakdown, feedback summary, or comparison data]

Constraints:

- Do not write correlation as causation
- Do not conclude success or failure unless there is a clear target to compare against
- Do not add cost, budget, or spend data that was not provided
- Label anything uncertain as "to be confirmed"

Output:

1. Event goals and data scope
2. Key observations
3. Stronger-performing areas
4. Weaker or anomalous areas
5. Possible causes (labeled as inferences)
6. Data still needed
7. Suggestions for next event

What to verify: Does AI know the event's time range and original goals? Did it write "correlated" as "caused"? Did it note missing cost, audience, or channel data?

8.10 System-exported CSV analysis

Situation: You have exported a CSV from a management system, form tool, survey platform, or data warehouse and need to understand the structure before drawing any conclusions.

Assignment:

Please review this system-exported CSV and summarize the fields, what analysis it can support, and any data quality issues.

Goal:

I do not need conclusions yet. I need to understand what this data can and cannot answer.

Data:

[Upload or paste the CSV header row and a sample of records]

Constraints:

- Do not draw overall conclusions from only the first few rows
- Do not overlook empty values, duplicate entries, or inconsistent formatting
- Do not modify the data
- If something cannot be determined from the data, say so explicitly

Output:

1. Field descriptions
2. Questions this data can answer
3. Questions this data cannot answer
4. Data quality issues found
5. Suggested analysis steps
6. Additional data that would be needed

What to verify: Did AI understand the fields before drawing conclusions? Did it surface empty values, duplicates, or formatting problems? Did it clearly separate "can analyze" from "cannot analyze"?

8.11 Requirements document into discussion items

Situation: A PM, designer, engineer, or operations person has received a requirements document and needs to organize it into open questions, scope, risks, and a priority for cross-team discussion.

Assignment:

Please organize the following requirements document into a cross-team discussion summary.

Goal:

I need to understand the purpose, feature scope, undefined questions, potential risks, and who to confirm next steps with.

Data:

[Paste the requirements document or meeting notes]

Constraints:

- Do not add features that are not in the document
- Do not estimate engineering effort
- Do not fill in unclear specifications as if they were confirmed
- Do not decide priority – only suggest a priority for discussion

Output:

1. Purpose of the requirement
2. User or department needs
3. Feature scope
4. Undefined questions
5. Risks and dependencies
6. Teams or people to confirm with
7. Suggested order for the next discussion

What to verify: Did AI add new requirements? Did it present suggestions as decisions? Are the open questions specific enough to be answerable?

8.12 Expense report and receipt review

Situation: A team member is preparing an expense report and wants to check completeness before submitting to finance.

Assignment:

Please review the following expense report data and check for missing items before it goes to finance.

Goal:

I need to know whether the documentation is complete, what is still missing, and how to write the purpose description. Do not approve the expense yourself.

Data:

[Paste expense dates, amounts, purposes, receipt types, and payment proof – mask sensitive financial or identification data]

Constraints:

- Do not judge whether the expense is reimbursable
- Do not ask for full credit card numbers, bank account numbers, or national ID numbers
- Do not make tax or approval decisions on behalf of finance
- List any missing documentation as "needs to be provided"

Output:

1. Expense summary
2. Documentation provided
3. Missing items
4. Suggested purpose description text
5. Items needing manager or finance confirmation
6. Pre-submission checklist

What to verify: Are amounts, dates, and currencies correct? Do receipts match payment records? Was no sensitive data printed in the output? Did AI avoid making the final approval decision?

8.13 Advanced case: admin console operation manual via Computer Use

This case demonstrates the full division of labor for an AI work assistant: the human controls login, permissions, and sensitive data; the AI work assistant uses Computer Use to observe the admin console, survey the page structure, capture safe screenshots, and draft an operation manual; the human then verifies and fills gaps.

This is an advanced scenario — not a beginner starting point. Public readers who want to try this should use only demo, training, staging, or read-only accounts. Do not use real customer data, production admin credentials, high-privilege accounts, or any system where the AI could modify live data.

Why this task suits AI

Many admin panels, member systems, order systems, content management systems, and internal tools contain dozens of pages, menus, buttons, and fields that need to be documented into a manual team members can actually use. Doing it manually means: open each page, take a screenshot, mask sensitive data, write the steps, format everything, and convert to a shareable document.

This is exactly the kind of work where AI earns its place: high volume of observation, description, and formatting — with the human retaining control over everything that matters.

What the human handles before handing off

AI must not handle login, credentials, or permission decisions. A safe setup:

1. Log in to the admin system yourself
2. Use a test, demo, staging, or read-only account whenever possible
3. Open the system to the starting page you want AI to observe
4. Decide which sidebar pages, tabs, and sections AI can navigate
5. Avoid pages that should not be exposed: payments, personal data, permissions, system settings
6. Explicitly tell AI it may only observe, navigate, screenshot, and organize — no modifications
7. Instruct AI to pause and ask if it encounters anything with sensitive data or uncertain function

Assignment

Please help me draft an operation manual for this admin console using Computer Use.

Goal:

Produce a manual team members can follow, covering main pages, functions, buttons, fields, and basic operating steps.

Data:

I am already logged in and have opened the system to the starting page you can view. If this is not a test, demo, staging, or read-only account, please remind me to confirm the permission level before proceeding.

You may:

- Observe menus, buttons, fields, and labels on the current page
- Click sidebar items, tabs, and expandable menus that only navigate between views
- Survey the main features and page structure
- Capture screenshots within the agreed safe scope and organize steps
- Produce a Markdown operation manual draft

You must not:

- Log in or request passwords

- Click add, edit, delete, save, submit, publish, modify settings, process payments, or change permissions
- Guess the purpose of any button you are not certain about – stop and ask
- Download or share data externally
- Screenshot any page showing customer names, phone numbers, transactions, payments, permissions, or credentials without first alerting me to mask or skip that screen

Output:

1. Feature overview
2. Sidebar page list
3. Purpose of each page
4. Main fields and buttons per page
5. Basic operating steps
6. Screenshot suggestions
7. Open questions
8. Items AI should not judge or operate

What the human must verify

- Did AI miss any sidebar pages or navigation sections?
- Does each step match the actual screen layout?
- Are screenshots masked for sensitive data?
- Did AI stay within observation and safe navigation — no write actions?
- Are unclear functions labeled as "to be confirmed"?
- Can a new team member actually follow this manual?

Division of labor

Stage	Human	AI work assistant
Before starting	Log in, open starting page, set allowed scope, declare off-limits areas	Wait for clear assignment
During	Control sensitive data and high-risk operations	Use Computer Use to observe pages, navigate safely, survey, screenshot, and organize
After first draft	Verify, flag missed pages, confirm screenshot safety	Revise, fill gaps, flag open questions
Final	Decide whether to deliver or save as a template	Produce a shareable draft

If AI misses pages in the first draft, that is not a failure — that is exactly what verification is for. Flag the missing sections and ask for a second draft.

8.14 Working methods from practice

The most valuable pattern when working with a Chat or workspace agent is not "ask one question, get one answer." It is building a consistent rhythm. These six methods apply across all scenarios in this chapter.

Method 1: Establish the source of truth first

Rather than:

What do you think should change in this document?

Try:

Please read the current document, the relevant background, and the actual data. Then assess what should change.
Separate confirmed facts, reasonable inferences, and items still to confirm.

This applies to: document revisions, SOP updates, customer complaint triage, project status reports, API or system behavior analysis.

Method 2: Define the deliverable before asking AI to execute

Rather than:

Please organize this.

Try:

Please organize this into a guide a team member can use from day one.
Include: quick start, full walkthrough, practice exercises, verification checklist, safety boundaries, and reusable template.

The clearer AI's picture of the final deliverable, the closer the first draft is to usable.

Method 3: For high-risk tasks, draw the boundary first

Whenever the task touches customers, payments, contracts, personal data, permissions, or any official submission or data modification, state the boundary upfront:

You may organize, check, and produce a draft.
You must not submit, delete, or modify live data, and must not make the final decision for me.
When in doubt, label it as "to be confirmed."

This is not about distrust. It is about making responsibility explicit.

Method 4: Ask for a retrospective, not just a result

Many useful Skills do not come from the first prompt — they come from reviewing a successful session afterward:

Please summarize this session: describe your understanding of the goal, the steps you took, the issues I pointed out, the rules we ended up with, the verification method, and the stopping conditions.

This turns one successful collaboration into a reusable process.

Method 5: Save the method, not just the output

When something will happen again, save how you did it — not just what AI produced:

- How did you assign the task?
- What data did you provide?
- Where did AI go wrong?

- How did you verify?
- What constraints should be added next time?
- Is this worth formalizing as a Skill?

This is the difference between individual use and organizational learning.

Method 6: Before trying a new tool, rewrite the assignment

When you see a new AI product or someone asks "is this tool better?" — do not compare features first. Rewrite your need as a clear assignment and try it with the tool you already have.

For example, instead of:

Is there a better AI tool for summarizing event reports?

Try:

I have an event results report with registration numbers, actual attendance, channel breakdown, a feedback summary, and a before/after comparison.

Please produce:

1. Key results
2. Anomalies
3. Possible causes (labeled as inferences)
4. Data still needed
5. Suggestions for the next event

Constraints:

- Do not write correlation as causation
- Do not add cost or budget data that was not provided
- Label anything uncertain as "to be confirmed"

If the tool you already use produces a verifiable first draft from this, you probably do not need a new tool. If it consistently falls short, then investigate why: is the data messy? Are fields missing? Is the task goal unclear? Does this genuinely require a different system connection? A real gap justifies switching; habit or novelty does not.

8.15 Advanced case: safe disk space cleanup

Many people's first response when their computer runs low on storage is to search for free cleanup software. This is a significant risk.

Some tools that appear to be legitimate disk cleaners may ask you to download a .dmg file, install a configuration profile, enter your system password, grant full disk access, and then read your browser data, email, cloud files, or saved credentials. The worst outcome is not a cluttered disk — it is having your account sessions, cloud data, or work system permissions stolen.

When you run low on storage, do not reach for cleanup software. A better approach: within the tools and environment your organization permits, ask a workspace agent such as a CLI agent to use built-in system commands to conduct a read-only survey first. Let AI map where the space is going, categorize candidates, and flag what must not be touched. Do not install tools from unknown sources, and do not run scripts you cannot read.

This is an advanced scenario, not a starting exercise. The public edition demonstrates only the survey, classification, and risk-labeling steps. Beginners should not let AI delete files directly. Before anything is moved to the trash, back up the data, confirm the path, confirm it is not a cloud-synced or work file, and make the final decision yourself.

What a workspace agent can do

- Report current disk usage and remaining space
- Identify the largest folders and file types
- Distinguish frequently used, likely unused, and clearly temporary or installer files
- Produce a cleanup candidate list with risk labels for each item
- Flag folders and types that must not be touched

What AI must not do

- Download or install any external cleanup software
- Delete files without a confirmed candidate list
- Touch browser data, email, credentials, cloud-synced folders, or work project files
- Use permanent deletion commands (bypassing the trash)
- Treat "not recently modified" as equivalent to "safe to delete"
- Assume any file is unimportant based on size alone

Assignment

Please help me safely assess a low disk space situation.

Constraints:

- Do not delete any files
- Do not download or install any cleanup software
- Do not touch passwords, browser data, email, cloud-synced folders, work projects, or system settings
- Do not use permanent deletion
- Begin with a read-only survey only

Please report:

1. Current disk space remaining
2. Largest folders or file types
3. Low-risk candidates you believe could be cleaned up
4. Candidates that need my confirmation before touching
5. High-risk folders or types that must not be touched
6. Estimated maximum space that could be freed

Please format the output as a table with columns:

path | size | type | likely use | frequency (frequent / infrequent / unknown)
| suggested action | deletion risk | confirmation needed

How AI should classify files

"Frequent" and "infrequent" cannot be determined by intuition alone. Useful signals:

- **File location:** Old installers in Downloads are typically lower risk than working files in Documents.

- **File type:** Old .dmg, .zip, exported files, and build artifacts are typically lower risk than project source files.
- **File size:** Prioritize large files — the efficiency gain is higher.
- **Last modified time:** A useful signal but never the sole reason to delete.
- **Reproducibility:** Caches, compiled outputs, and re-downloadable exports are safer than originals.
- **Work context:** Active projects, client data, expense documentation, and asset libraries must not be deleted just because they are large.

Items typically safe to review first

- Installer files and archives in Downloads that have not been opened in a long time
- Screen recordings, meeting recordings, or temporary exports that are clearly no longer needed
- Caches and build artifacts for projects you are not currently working on
- Duplicate exports, CSVs, PDFs, or reports you have already submitted
- Delivered and backed-up assets that can be re-downloaded if needed

Even for low-risk items, produce a list for human confirmation first.

Items to leave alone

Beginners should not allow AI to clean up:

- Application support data in system library folders from unfamiliar apps
- Browser profiles, cookies, sessions, or saved credentials
- Email, messages, keychain, or system configuration
- Cloud-synced folders (iCloud, Google Drive, Dropbox, etc.)
- Work projects, client data, financial records, contracts, or expense originals
- Photo libraries, music libraries, or video editing project source files
- Hidden folders whose purpose is not clearly understood

These areas are not permanently off-limits — they just must not be cleaned up by AI without explicit context and human confirmation at each step.

After confirming the candidate list

Please help me plan the cleanup steps for the low-risk items I confirmed.

Constraints:

- Do not delete anything yet
- For each item, describe what to do, where it would go, and any possible side effect
- Do not act on any path outside the confirmed list
- If a path does not exist or the content differs from the earlier survey, stop and report before proceeding
- I will confirm each item individually before moving it to the trash

Output:

1. Items suggested for trash
2. Size of each item

3. Reason for the suggestion
4. Possible impact
5. Items that need manual review or should be kept

What this scenario is really teaching

When your disk is full, do not download unknown cleanup software. Use a workspace agent within your organization's permitted tools to do a safe read-only survey first. Then decide yourself what to remove.

AI's role here is to help you see where space is going, sort cleanup candidates by risk level, and protect you from touching high-risk data by accident. The person who confirms and acts is always you.

Practice checklist:

- Choose one scenario from this chapter that matches work you do regularly.
- Write down: the data source, the constraints, and how you will verify AI's output.
- Check: did you confirm permissions and safety scope before letting AI access any system or local file?

9. Diagnosing Bad Assignments

What you will learn:

- Understand why vague or risky assignments produce inconsistent results.
- Learn to rewrite assignments that are too short, too ambiguous, or too high-risk.
- Build the ability to diagnose problems in your own prompts before you run them.

9.1 Too short, too vague

Help me organize this.

Why this fails — AI has no way to answer these questions:

- Who is this for?
- How much detail is needed?
- Should it preserve every nuance or just the highlights?
- What format should the output take?
- Is it allowed to fill in gaps with inferences?

When an assignment leaves all these decisions to the AI, the results will differ every time you run it. You get lucky sometimes, but you cannot rely on it.

9.2 Good structure: goal, format, constraints

Please organize the following customer conversation into an internal handling summary.

Purpose: Let the customer service manager quickly understand the current status of this case.

Output:

1. Customer's request
2. Background
3. What we have committed to
4. What is still unconfirmed
5. Possible risks
6. Suggested next step

Constraints:

- Do not add information not in the conversation
- Do not decide whether to refund
- Label anything uncertain as "to be confirmed"

This works because it answers all five questions the vague version left open. The AI knows who reads it (the manager), what shape it takes (six numbered sections), and two hard boundaries it must not cross.

9.3 Letting AI decide high-risk items

This customer is angry. Tell him we can give a refund.

Why this is dangerous: Refund decisions belong to a manager or finance team. AI cannot commit on behalf of your company, and an AI-generated promise may create a liability your company did not intend to make.

Better — gather information first, commit to nothing:

This customer is upset. Please help me organize:

1. Why the customer is upset
2. What the customer is asking for
3. What we have previously committed to
4. Whether there is any refund-related information
5. Questions we need manager confirmation on before replying

Do not write a formal reply yet.

This gives you the information you need to make a responsible decision, without handing that decision to the AI.

Practice & self-check:

- Find one vague assignment you have written before and rewrite it with a clear goal, output format, and at least one constraint.
- When AI output varies between runs, diagnose whether the cause is missing data, unclear goal, or an undefined risk boundary.
- Can you identify a problem in an assignment before you run it?

10. Verification: Do Not Accept Output Blindly

What you will learn:

- Apply four verification questions to every AI output before using it.
- Recognize the four most common ways AI output goes wrong.
- Build the habit: reviewing output matters more than writing a perfect prompt.

10.1 Four verification questions

Every time you receive AI output, run these four checks before accepting it:

1. **Did it miss anything?** Is important information from the source not reflected in the output?
2. **Did it guess?** Did it turn "possibly" into "certainly" — presenting an inference as a fact?
3. **Did it overstate?** Did it commit to something that no one has actually authorized?
4. **Can it be handed off?** If someone else reads this, can they take the next action without coming back to ask you?

These four questions apply regardless of how good the prompt was. Even a well-structured assignment can produce output with one or more of these problems.

10.2 Common errors

Error 1: Inference stated as fact

AI writes: **"The customer is likely dissatisfied because of the price."**

If the original conversation did not mention price, this is a guess dressed up as analysis.

Correct version: **"Reason for dissatisfaction not clearly stated; may relate to price but needs confirmation."**

Error 2: Added information

AI writes: **"We will complete the process within three business days."**

If your company has not committed to this timeframe, it cannot appear in the output — especially in anything that might reach the customer.

Correct version: **"We will confirm the case status and reply with the next steps as soon as possible."**

Error 3: Advertising tone

AI writes: **"We deeply value your precious feedback and will strive to provide the most outstanding service."**

This sounds polished but says nothing useful and can come across as dismissive.

Replace with: **"We have received your feedback and will confirm the situation, then reply with the available next steps."**

Error 4: Output looks complete but cannot be executed

AI writes: *"We recommend improving internal processes, strengthening communication, and increasing efficiency."*

This is too abstract. There is no owner, no next step, and no way to know when it is done.

Ask AI to revise: *"Please rewrite each recommendation as a concrete action item with: what to do, who is responsible, what data is needed, and how to know it is done."*

Practice & self-check:

- Take one piece of recent AI output and run the four verification questions against it.
- Write the correction instructions for any errors you find.
- Can you clearly separate what AI confirmed from what it inferred?

11. Refinement: Make the Second Version Useful

What you will learn:

- Give AI specific, targeted feedback instead of just saying "redo this."
- Use the refinement formula to produce a more useful second version.
- Compare before and after to see exactly what changed.

11.1 Do not just say "rewrite"

Poor feedback:

Wrong. Rewrite.

AI does not know what was wrong. Without knowing which part to fix, it will produce another version with the same problems, or fix one thing and introduce a new one.

Better feedback:

This version has problems:

1. The tone sounds like advertising copy
2. It added a commitment that was not in the original
3. It did not list the open questions

Please rewrite it as an internal customer service summary. Keep the customer's original request. Do not add a handling timeline.

Now AI knows exactly what the problems are, what the target format is, and specific things it must and must not do.

11.2 Refinement formula

Use this structure every time you send a correction:

This version needs revision.

Problems:

- [Problem 1]
- [Problem 2]

Reason:

- [Why this is a problem for this particular use case]

Please revise to:

- [What you want instead]

Must keep:

- [What should stay unchanged]

Must not include:

- [What should be removed]

11.3 Worked example

Original AI output:

We sincerely apologize for the inconvenience caused. We will immediately process your refund.

Problems with this output:

- "Immediately process your refund" is an unauthorized commitment — refunds require internal review.
- The tone is generic and does not acknowledge the specific situation.

Refinement instruction:

This version cannot include "immediately process your refund" because refunds require internal confirmation first.

Please revise to:

- Acknowledge receipt of the customer's concern
- State that we will verify the order status
- Do not commit to a refund
- Do not commit to a timeline

Give me one version suitable for an initial customer reply.

Revised output direction:

Thank you for reaching out. We have received your concern and will look into the order status. We will follow up with you once we have confirmed the available next steps.

This version is honest, professional, and does not make a promise the company has not authorized.

Practice & self-check:

- Convert a "rewrite" instruction into a structured refinement using the formula above.
- Compare the first and second versions — identify exactly what improved.
- Did you point to specific problems rather than just expressing general dissatisfaction?

12. Safety Boundaries

What you will learn:

- Establish clear data and action boundaries that apply every time you use an AI work assistant.
- Learn to de-identify sensitive data before sending it.
- Know which types of content always require human sign-off before going out.

Chapter 3 covered what to avoid when you are just getting started. This chapter focuses on the boundaries that apply every time, at every stage.

12.1 Data that should not be pasted into AI without authorization

Unless your organization has a clear authorization and a secure AI environment in place, do not paste:

- Passwords, API keys, or access tokens
- National ID numbers or passport numbers
- Credit card or bank account numbers
- Full customer personal records (name + contact + transaction combined)
- Unpublished financial data or internal reports
- Unsigned or confidential contracts
- Internal strategy documents

If you are unsure whether your organization has authorized AI use for a given type of data, treat it as not authorized until you confirm.

12.2 How to handle sensitive data: de-identify first

Before sending sensitive data to AI, remove or replace the identifying details.

Original:

John Smith, phone 555-0142, order #A12345, requesting a refund of \$1,280.

De-identified version:

Customer A, one order, requesting a refund (amount omitted).

Assignment to AI:

The following data has been de-identified. Please help organize:

1. Case background
2. Customer's request
3. Confirmed information
4. Open questions that need follow-up

Do not attempt to infer or reconstruct the customer's identity.

12.3 De-identification reference

Original	De-identified version
Full customer name	Customer A
Phone number	Phone number masked
Order number	Order number masked
Specific dollar amount	Approximate amount, or omitted
Internal staff name	Team member A / Manager B
Credit card number	Last four digits only
Bank account number	Last five digits only

The goal is to preserve enough context for AI to be useful while removing anything that could expose an individual's identity.

12.4 External content always requires human confirmation before sending

AI can draft any of the following, but a human must review and approve before it is sent:

- Formal customer replies
- External announcements
- Social media posts
- Contract terms
- Price quotes
- Apology statements
- Recruitment or HR notices
- Financial-related communications

Drafting and sending are two separate steps. Never compress them into one.

Practice & self-check:

- Take a sample piece of data that contains sensitive information and produce a de-identified version.
- List three types of data from your own work that should not be pasted into AI directly.
- Do you know which situations require checking your organization's policy before proceeding?

13. Assignment Quality Levels

What you will learn:

- Use a five-level scale to assess the quality of your assignments.
- See how adding goal, format, constraints, and verification standards each make output more stable and usable.
- Upgrade a one-line assignment into a handoff-ready task.

Use this scale to check your assignments before you send them.

Level 1: One-line assignment

Help me organize this.

Problem: Too vague. AI does not know the purpose, the audience, or the expected format. Every run will produce different output.

Level 2: Has a goal

Help me organize this into a summary for my manager.

Better than Level 1, but still no format. AI will guess what a "summary" means, and its guess may not match yours.

Level 3: Has a format

Help me organize this into a summary for my manager, divided into background, current status, risks, and next steps.

Can produce a usable first draft. The four sections give AI a clear structure to fill, which reduces variation between runs.

Level 4: Has constraints

Help me organize this into a summary for my manager, divided into background, current status, risks, and next steps. Do not add information not in the source; label anything uncertain as "to be confirmed."

Reduces guessing and fabrication. The two constraints tell AI what it must not do, which is often as important as telling it what to do.

Level 5: Has a verification standard

Help me organize this into a summary for my manager, divided into background, current status, risks, and next steps. Do not add information not in the source; label anything uncertain as "to be confirmed." Every next step must include a responsible role, what data is needed, and a completion standard.

Closest to a handoff-ready, verifiable assignment. Someone else can read this output, know what is done and what is open, and take action without coming back to you for clarification.

Practice & self-check:

- Take a Level 1 assignment from your own work and upgrade it step by step through all five levels.
- After each level, note what AI now knows that it did not know before.
- Does your assignment include a verification standard?

14. Confidence Labels

What you will learn:

- Ask AI to separate confirmed facts, reasonable inferences, and items still to be verified.
- Reduce the risk of AI presenting guesses as facts.
- Build this habit especially when working with incomplete data.

AI does not naturally flag its own uncertainty. Without instruction, it will state inferences and confirmed facts in the same tone. Asking for confidence labels forces the separation.

14.1 Basic template

Add this to any assignment where data completeness matters:

Please divide the output into:

1. Confirmed information
2. Reasonable inferences
3. Insufficient data - needs human confirmation

14.2 When to use it

This is especially valuable for:

- Customer complaint summaries
- Report anomalies or discrepancies
- Manager decision briefs
- External reply drafts
- Expense or reimbursement checks
- Any case where the input data is incomplete or inconsistent

If you are about to share or act on AI output, ask yourself: do I know which parts of this are confirmed and which are guesses? If not, use confidence labels.

14.3 Example output

Confirmed information:

- Customer emailed on May 1 saying the product was not received.
- Order status currently shows "shipped."

Reasonable inferences:

- Possibly a logistics delay or incorrect delivery address, but not yet confirmed.

Insufficient data - needs human confirmation:

- Logistics delivery record
- Whether the delivery address on file is correct
- Whether there is an existing customer service thread for this order

This structure makes it immediately clear to any reader — including your manager — exactly what is known, what is a working hypothesis, and what needs to be looked up before a decision can be made.

Practice & self-check:

- Pick a case with incomplete data and ask AI to output it using the three-tier confidence structure.
- Check: did AI label its inferences as inferences, or state them as facts?
- Can you accept "insufficient data" as a valid and useful output — not a failure?

15. Pre-Send Checklist for External Replies

What you will learn:

- Build a risk check routine before sending any external content.
- Know what categories of risk to look for.
- Learn to use AI as a checker, not a sender.

Sending an AI-drafted message without reviewing it is one of the most common ways AI use goes wrong in practice. This checklist is your last gate before anything leaves your organization.

15.1 Checklist

Before sending any external content, verify each of the following:

- **No unauthorized commitments** — no promise of a specific price, refund, compensation, discount, or handling timeline unless already confirmed internally
- **No admission of liability** — no language that explicitly or implicitly accepts fault on behalf of your organization
- **No internal information leaked** — no internal process details, staff names, internal codes, or system names that should not be shared externally
- **No inappropriate identifiers** — no customer data belonging to a different case, no internal reference numbers that expose system structure
- **Tone is stable and non-escalating** — not defensive, not dismissive, not condescending
- **Inferences are not presented as facts** — every uncertain point is labeled as pending or to be confirmed
- **Escalation reviewed** — confirmed whether this needs sign-off from manager, finance, legal, or PR before sending
- **Consistent with existing policy** — aligns with known company rules, past commitments, and standard handling
- **Leaves room for follow-up** — does not close off options or make a final-sounding statement when the situation is still open

15.2 Use AI to check before sending

Please check this external reply draft for the following risks:

1. Improper commitments
2. Admission of liability
3. Leaked internal information
4. Escalating language
5. Inference presented as fact
6. Items that need manager confirmation but are not labeled

Please only list the problems and suggested changes. Do not rewrite the entire message.

The key phrase here is "do not rewrite the entire message." You want a list of specific issues, not a replacement — that way the final decision stays with you.

Practice & self-check:

- Find a past external reply and audit it against this checklist.
- Ask AI to list only the risks, not rewrite the content.
- Is the final human review step still in place before anything is sent?

16. Template Library

What you will learn:

- Get ready-to-adapt templates for the most common workplace tasks.
- Understand why templates must be personalized — not used as-is.
- Build your own personal library over time.

Every template in this chapter needs you to fill in your specific goal, source data, and constraints before using. Pasting a template directly without adapting it is the most common mistake new users make.

16.1 Meeting notes → action items

Please organize the following meeting notes into an action-item tracking table.

Output:

1. Decisions made
2. Action items
3. Owner (person or team)
4. Deadline
5. Risks or dependencies
6. Open questions still to be resolved

If something was not mentioned in the notes, write "not specified." Do not fill in gaps.

16.2 Customer message → internal summary

Please organize the following customer message into an internal summary.

Output:

1. Customer's request
2. Customer's apparent emotional state
3. Confirmed facts
4. What we have already committed to
5. What we cannot commit to at this stage
6. Suggested next step
7. Open questions for follow-up

16.3 Customer reply → polished version

Please rewrite the following into a version suitable to send to the customer.

Tone: Clear, polite, non-confrontational, non-overcommitting.

Constraints: Keep the original meaning. Do not add any new commitments. Do not mention internal processes.

16.4 Document draft → logic check

Please review the following document draft.

Find:

1. Logic contradictions
2. Missing information
3. Tone that is inappropriate for the intended audience

4. Sections that are easily misunderstood
5. Suggested revision directions

Do not rewrite. Just identify the issues and recommend directions.

16.5 SOP draft

Please organize the following process into a standard operating procedure (SOP).

Include:

1. Applicable situations
2. What to prepare before starting
3. Step-by-step operating instructions
4. Notes or cautions for each step
5. Common mistakes and how to avoid them
6. Completion checklist
7. When to escalate to a manager

16.6 Report anomaly summary

Please review the following report data and identify anomalies.

Separate your output into:

1. Confirmed anomalies
2. Possible causes (label these as inferences, not conclusions)
3. Data still needed to verify
4. Priority recommendation for follow-up
5. Items you are uncertain about

Do not state inferences as conclusions.

16.7 Manager decision brief

Please organize the following into a decision brief for manager review.

Output:

1. Background
2. Current status
3. Available options
4. Risks of each option
5. Recommended option with reasoning
6. What the manager needs to decide before we proceed

16.8 Activity planning breakdown

Please break the following activity goal into an execution plan.

Output:

1. Phase breakdown
2. Tasks within each phase
3. Suggested sequencing
4. Who needs to be involved
5. Key risks
6. What can be done in the first week

16.9 FAQ generation

Please organize the following source material into a FAQ.

For each item, output:

- Question
- Standard answer
- What we cannot commit to in the answer
- When to escalate to a human

Do not add policies or information not present in the source material.

Practice & self-check:

- Pick one template from this library and adapt it to a real task you have coming up.
- Fill in the specific goal, data source, constraints, and verification approach.
- Did you avoid pasting the template directly without customizing it?

17. Red-Flag Rules

What you will learn:

- Know exactly when to stop and ask a human instead of continuing.
- Distinguish between data gaps, high-risk decisions, formal operations, and permission boundaries.
- Reframe stopping as risk control — not as failure.

The most important skill in working with an AI work assistant is not knowing when to go faster — it is knowing when to stop.

17.1 Stop and ask a human when:

- **Data is insufficient but a decision is required** — do not let AI fill in the gaps when the answer actually matters
- **The task involves money, refunds, contracts, or personal data** — these require a human decision chain, not AI judgment
- **An external commitment is about to be made** — any statement to a customer, partner, or regulator must be human-reviewed first
- **A customer threatens a complaint, lawsuit, or public disclosure** — this escalates the stakes beyond routine handling
- **You cannot understand what AI produced** — if you cannot verify it, you cannot be responsible for it
- **The same problem persists through three revision attempts** — repeated failure is a signal the task needs human input, not another prompt
- **The output affects other departments** — cross-team impact requires cross-team sign-off
- **The action involves deleting, overwriting, publishing, or changing permissions** — these are hard to reverse; confirm before proceeding
- **The task requires formal legal, tax, HR, or financial judgment** — AI can help organize the information, but the judgment call must be human

17.2 Why stopping is the right move

When you stop and ask, you are not admitting that AI failed. You are recognizing that the task has reached a boundary where the cost of being wrong is higher than the cost of slowing down.

An AI work assistant is a fast, capable tool for organizing information and producing first drafts. It is not a decision-maker, an authority, or a substitute for organizational process. Keeping that distinction clear is what makes you a responsible and effective user.

Practice & self-check:

- Recall one task where you should have stopped earlier than you did. What was the signal you missed?
- Categorize it: data gap, permission boundary, high-risk decision, or formal judgment?
- When you stop, do you treat it as risk control rather than failure?

Farceur Liu | farceurliu.github.io/ai-butler-handbook/

FREE PUBLIC EDITION | Farceur Liu | farceurliu.github.io/ai-butler-handbook/

FREE PUBLIC EDITION | Farceur Liu | farceurliu.github.io/ai-butler-handbook/

FREE PUBLIC EDITION | Farceur Liu | farceurliu.github.io/ai-butler-handbook/

FREE PUBLIC EDITION | Farceur Liu | farceurliu.github.io/ai-butler-handbook/

FREE PUBLIC EDITION | Farceur Liu | farceurliu.github.io/ai-butler-handbook/

FREE PUBLIC EDITION | Farceur Liu | farceurliu.github.io/ai-butler-handbook/

FREE PUBLIC EDITION | Farceur Liu | farceurliu.github.io/ai-butler-handbook/

FREE PUBLIC EDITION | Farceur Liu | farceurliu.github.io/ai-butler-handbook/

FREE PUBLIC EDITION | Farceur Liu | farceurliu.github.io/ai-butler-handbook/

18. Turning Reusable Workflows into Skills

What you will learn: Understand that a skill is a saved, verified workflow — not just a polished prompt. Learn the five-step path from collaboration to skill packaging. Know when a template is enough and when a formal skill is worth building.

18.1 What is a skill

A skill is a set of standing instructions written for your AI assistant — not a single saved prompt, but a complete workflow that has been run, corrected, and verified, then preserved so it can be reused consistently.

Think of it as: *the standard operating procedure you wrote for your AI work assistant.*

Story: how a skill grows from repeated work

Maya reconciles expense reports every month. The first time she tried AI assistance, she wrote:

Help me organize these receipts.

The AI produced a tidy list of dates and amounts — but missed that some receipts were missing payment confirmation, and it said nothing about vague expense descriptions.

She refined her approach the next month:

Please organize this expense data.
Do not decide whether expenses are approvable.
List: what documentation is provided, what is missing,
which descriptions need more detail, and a pre-submission
checklist to run before sending to finance.

That worked much better. She noticed the same problems appeared every month: missing invoices, missing payment proof, descriptions too short, amounts that did not match screenshots.

So she asked AI to review what had just worked and distill it into a repeatable process. That became the draft for her "Expense Data Preparation and Missing-Document Check" skill.

The point of the story: **skills are not written from scratch at the start. They grow out of repeated work.**

18.2 When is a workflow worth turning into a skill

A workflow is worth packaging as a skill when it:

- Happens every week
- Is done by multiple people
- Requires re-explaining the rules every time
- Has frequent errors
- Accepts consistent input types
- Produces consistent output format
- Has a clear verification method

If fewer than half of these apply, keep using a prompt template for now.

18.3 What a good skill includes

These are the fields worth filling in. A complete, copyable template is in section 18.7.

Field	Purpose
Skill name	Tells anyone immediately what work this handles
Applicable situations	When to use it; when not to
Required data	What to prepare before running
Operating steps	What AI should do, in what order
Output format	What the final output looks like
Verification method	How the human checks whether the result is usable
Common errors	What AI typically misses, guesses, or overstates
Stop conditions	When to pause and ask a human instead
Helper files	Advanced skills may include .sh or .py scripts, templates, or examples

18.4 Why large workflows need incremental packaging

Do not start by asking AI to "create a skill for me." Complex workflows contain multiple phases, different types of data, various roles, and edge cases that are not obvious until you have run the process several times.

A skill produced in one shot looks complete but is built on guesses. The reliable path is:

Step	What to do	Why
Collaborate first	Run the actual work step by step with AI	AI learns the real process, not an imagined one
Verify each phase	Check errors and add constraints as you go	Discover where AI tends to guess or misjudge
Debrief after completion	Ask AI to summarize what worked and where human judgment was needed	Turn the interaction into a discussable process
Evaluate	Assess whether this is ready to be a skill	Avoid packaging an immature process
Package last	Only then write it as a template or skill draft	Built on verified behavior, not speculation

To start a large workflow collaboratively:

This is a larger workflow. Please do not rush to create a skill.

Please help me work through this step by step.

After each phase, report:

1. What you did
2. What data you used to decide
3. What you were uncertain about
4. What needs my verification

5. What you recommend for the next phase

After we finish the whole task, we can review whether it makes sense to package this as a template or skill.

18.5 What tools can generate a skill

Two things are worth separating:

- **Generating a skill draft** — producing a document that describes the workflow for AI to follow.
- **Installing a usable skill** — placing that draft in the right location and format so a workspace tool actually reads it.

These require different tools.

What you want to do	Tools that work	Notes
Draft a skill from scratch	Any chat AI (ChatGPT, Claude, Gemini, etc.)	Any conversational tool can help you draft
Retrospect a workflow into a skill	Any chat AI	Summarize what just worked into a reusable format
Edit an existing skill Markdown file	Workspace tools (Claude Code, Codex, Gemini CLI)	Need to read, edit, and check files
Create a tool-readable SKILL.md	Workspace tools or manual editing	Need to know the correct folder and format
Add a .sh or .py helper script	Workspace tools, or someone who can review scripts	Scripts execute real actions; safety review is required
Install third-party skills or packages	Not recommended for beginners	Requires security review, source verification, and license checks

Scripts are not required for a beginner skill. Most people start with a clean SKILL.md and never need a helper script.

What a general chat AI can do

ChatGPT, Claude, Gemini, and similar tools can help you:

- Draft a skill
- Rewrite tone or language
- Fill in missing steps
- Add stop conditions
- Add verification methods
- Turn a successful interaction into a reusable template

What they typically cannot do: place the file in your team's designated folder, or confirm whether your local tool will actually load the skill.

A skill draft from a general chat AI is more like:

Content ready, but not yet installed or deployed.

What a workspace-type AI assistant can do

Claude Code, Codex, Gemini CLI, and similar tools are better suited for deploying a skill:

- Read your current folder or repository
- Find existing skill formats or documentation
- Create or modify Markdown files
- Create or adjust .sh / .py helper scripts
- Check section structure
- Verify that scripts stay within allowed scope
- Report exactly which files were changed
- Run allowed check commands

If you want to build a skill that your tools will actually load, delegate it like this:

Please do not modify any files yet.

First, check whether the current project or folder contains any existing skill format, examples, or documentation.

Goal: I want to package the following verified workflow as a reusable skill: _____

Please report back:

1. Which relevant files you found
2. Where skills should be placed
3. What filename or format is required
4. What information is still missing from the skill draft
5. What needs my confirmation before you proceed

Wait for my confirmation before creating or modifying files.

After confirming, follow up with:

Please create or update the skill file using the format we just agreed on.

Constraints:

- Only modify files related to this skill
- Do not install external packages
- Do not download external content
- Do not change unrelated settings
- If the file location is unclear, stop and ask me

When done, report:

1. Which files were changed
2. The skill's applicable situations
3. Operating steps
4. Safety constraints and stop conditions
5. What still needs manual review

When a draft is enough

If the skill will be read by teammates — copied, referenced, adapted — storing it in a shared document or knowledge base is fine. It does not need to be a tool-readable SKILL.md.

Examples that work well as team templates without becoming formal skills:

- Customer service response checklist
- Meeting notes organization method
- Expense documentation checklist
- Event results summary format

Start as a team template. Once it is stable, used repeatedly by multiple people, and has clear stop conditions, then consider whether a formal skill is worth building.

18.6 Five steps: from first run to skill draft

The five steps from section 18.4 look simple in a table. Here is how each step works in practice.

Step 1: Run the real task once

Pick something you actually repeat, for example:

- Organizing weekly customer support cases
- Reconciling monthly expense reports
- Summarizing results after every event
- Capturing action items after every meeting
- Listing open questions after receiving a requirements document

Run it once with a plain prompt. See whether AI can genuinely help.

Step 2: Record how you delegated

Preserve three things:

1. What raw input did I give the AI?
2. How did I phrase the delegation?
3. Which parts of the output were useful, and which needed correction?

This matters because skills are not invented from imagination — they are extracted from work that has actually run.

Step 3: Distill the fixed rules

Ask yourself:

- What data needs to be ready before this task starts?
- What should AI always do first?
- What format should the output be in?
- What must AI never do?
- Under what circumstances should it stop and ask a human?

- How will you check the result?

If you cannot answer these questions yet, the workflow is not ready to become a skill.

Step 4: Ask AI to draft the skill

Please turn the following verified workflow into a skill draft.

Goal: this skill is standing instructions for an AI assistant.
When the same type of task arrives, the assistant follows this process instead of improvising.

When this task applies:

Input I usually provide:

What I want AI to do:

Output format I want:

What AI must not do:

When to stop and ask a human:

How I will verify the result:

Please output as:

1. Skill name
2. Purpose
3. Applicable situations
4. Required data
5. Operating steps
6. Output format
7. Safety constraints
8. Stop conditions
9. Verification method
10. Common errors

Step 5: Test on new cases and refine

A skill draft is not the final version. Test it on at least two or three different inputs:

- One normal case
- One case where data is incomplete
- One case with higher risk or a situation that should trigger a stop condition

After each test, ask:

- Did AI follow the steps?
- Was the output format consistent?
- Did it overpromise or make unauthorized judgments?
- Did stop conditions activate when they should?

- Was the result easy to verify?

Only promote to an official skill after it passes all three test types.

18.7 Skill draft template

Copy and fill in the blanks:

Skill Name

[Name here]

Purpose

This skill helps [describe use case].
Its goal is not [what it does not do], but [what it does].

Applicable Situations

Use when [description].

Suitable for:

- [situation]
- [situation]

Not suitable for:

- [situation]
- [situation]

Required Data

The user should provide:

- [item]
- [item]
- [item]

If data is missing, list what is missing and do not guess.

Operating Steps

1. Confirm the task goal and data scope
2. Organize confirmed information
3. Label reasonable inferences as inferences
4. List items to be confirmed
5. Output in the specified format
6. Check whether any stop conditions apply

Output Format

Please output:

1. [item]
2. [item]
3. [item]

Safety Constraints

AI must not:

- [action]
- [action]

Stop Conditions

If any of the following occur, stop and alert the user:

- [condition]
- [condition]

Verification Method

The user should check:

- [check]
- [check]

Common Errors

AI frequently makes these mistakes:

- [error]
- [error]

Correction: [how to fix]

Helper Files (Advanced)

Does this skill need .sh, .py, templates, or example files?

If no: keep only SKILL.md.

If yes, for each file describe:

- File name and purpose
- What it reads
- What it modifies or produces
- What it must not touch
- What needs human confirmation before running
- Whether it supports dry-run

18.8 Example: turning a meeting-notes prompt into a skill

Here is the same workflow at two levels of maturity.

Original one-time prompt

Summarize the following meeting notes into action items.
Include decisions, owners, deadlines, risks, and open questions.
Do not add anything that was not mentioned in the notes.

After packaging as a skill

Skill Name:

Meeting Notes Action-Item Digest

Purpose:

Help the user turn meeting notes into trackable internal action items.
This skill does not make decisions on behalf of the team,
and does not add items that were not raised in the meeting.

Applicable Situations:

User provides meeting notes, a transcript, rough notes, or a
post-meeting summary and needs them organized into decisions,
action items, and open questions.

Required Data:

- Meeting notes or transcript
- Meeting date
- Meeting topic
- Participating departments or roles (if available)

Operating Steps:

1. Identify the meeting topic and context
2. Extract decisions that were clearly made
3. Extract action items
4. Record owners and deadlines; mark "not specified" if not mentioned
5. List risks and open questions
6. Verify that no information was added beyond what the notes contain

Output Format:

1. Meeting summary
2. Decisions made
3. Action items
4. Owners
5. Deadlines
6. Risks
7. Open questions

Safety Constraints:

- Do not add decisions that were not in the notes
- Do not assign owners without explicit mention
- Do not commit to deadlines on behalf of the team
- Do not make prioritization decisions for the team

Stop Conditions:

- Notes are too incomplete to determine action items
- Content involves financial, legal, HR, or contractual decisions
- User asks AI to decide priorities on the team's behalf

Verification Method:

Check that no key decisions were missed, that owners and deadlines were not invented, and that each action item is clear enough for someone else to act on.

18.9 Pre-skill checklist

Before building a skill, confirm each item:

- ☐ This task has been run at least two or three times
- ☐ Input data type is consistent across runs
- ☐ Output format is consistent across runs
- ☐ Common AI errors are already known
- ☐ You know when to stop and ask a human
- ☐ There is a clear human verification method
- ☐ The skill will not allow AI to make decisions, send outputs, delete data, or make commitments autonomously

If three or more of these cannot be confirmed, keep testing with a template first.

18.10 Skills and helper scripts

A beginner skill is a single SKILL.md file. A more mature skill may grow into a small folder:

```
my-skill/  
  SKILL.md  
  scripts/  
    check-input.sh  
    parse-report.py  
  templates/  
    output-template.md  
  examples/  
    sample-input.csv  
    sample-output.md
```

Each file type plays a different role:

File type	Purpose
SKILL.md	Describes when and how to use the skill, and what is off-limits
.sh scripts	Local commands, file organization, environment checks, batch operations
.py scripts	Data parsing, format conversion, CSV/JSON processing, complex checks
templates/	Fixed output formats, document templates, reporting formats
examples/	Sample inputs and outputs, so AI knows what success looks like

Scripts are fundamentally different from text instructions. Text only describes. Scripts may actually read files, modify files, delete files, connect to the network, or execute commands.

Any skill containing .sh or .py files requires extra scrutiny before running:

- Which folders will the script read?
- Will it modify or delete files?
- Will it connect to the network or download anything?
- Will it read passwords, tokens, API keys, browser data, or personal data?
- Does it require a system password or elevated permissions?
- If it fails mid-run, does it leave broken files or corrupt the original data?
- Is there a dry-run or read-only mode to test before committing?

Rule for beginners:

You can ask AI to generate a script draft, but do not run scripts you do not understand. Ask AI to explain every section: what it does, which files it touches, what the risks are — then decide whether to run it.

Before asking AI to write a script, ask for an assessment first:

This skill may need a .sh or .py helper script.

Please do not create or run any scripts yet.

Please first evaluate:

1. Does this workflow genuinely need a script, or is SKILL.md alone enough?
2. If a script is needed, which repetitive step should it handle?
3. Which files or folders would the script read?
4. Would it modify, delete, download, upload, or connect to the network?
5. Can it support a dry-run or read-only mode first?
6. What human confirmation is needed before running it?
7. If no script is used, what is the alternative approach?

Report the risks and recommendations before doing anything.

18.11 More skills is not better

The purpose of skills is to stabilize mature workflows — not to accumulate a collection of prompts.

Cases where packaging as a skill is premature:

- Done only once
- Input data varies significantly every time
- Output format is not yet stable
- No clear way to verify the result
- Risk boundaries are not yet defined
- The real motivation is offloading judgment to AI

The healthy progression is:

One-time prompt → personal template → team template → skill

Master stable delegation before you think about packaging.

18.12 Reverse-engineering a skill from a successful interaction

In practice, the most reliable path to a good skill runs backward from a completed task:

1. Let AI read the source material
2. State clearly what needs to be delivered
3. Run a first attempt
4. Human points out problems
5. AI revises
6. Human confirms the result is usable
7. AI retrospects on what worked
8. Distill into a template or skill

This is more reliable than writing a skill at the start, because you are packaging what actually ran — not what AI imagined.

Note: completing a task and retrospecting does not automatically mean you should create a skill.

Some tasks, even when they worked, may rely too heavily on human judgment, vary too much from run to run, have unclear risk boundaries, or lack a stable process. These are better saved as case notes or personal templates first. A skill can come later, once the pattern is settled.

After completing a task, ask AI to evaluate suitability — not to package immediately:

Please do not write this up as a skill yet.

Please evaluate whether this task is suitable to become a skill.

Break down:

1. Parts of this task that were stable and repeatable
2. Parts that still required significant human judgment
3. Parts where input varied significantly
4. Parts where risk boundaries were unclear
5. If a skill is appropriate, which scope should it cover?
6. Is this better as a skill, a template, an SOP, or just a case record?

Do not create a skill just to create a skill.

First: the debrief prompt

After the task is done, start here:

Please summarize what just happened in this interaction.

Output:

1. What you understood the task goal to be
2. What data I provided
3. The steps you actually took
4. Problems I pointed out along the way
5. The judgment rules that ended up being used
6. Boundaries you respected – what you did not do or guess
7. How the human should verify the result
8. What questions to ask first if the same type of task arrives again

This step is not about generating a skill — it is about confirming AI understood the work. If the debrief is wrong, correct it before packaging anything.

Second: the packaging prompt

Once the debrief is confirmed, ask:

Please turn the workflow we just verified into a reusable skill draft.

Include:

1. Skill name
2. Trigger conditions
3. Data the user needs to provide
4. AI's operating steps
5. Output format
6. Verification method
7. Common errors
8. Safety constraints
9. Stop conditions

Constraints:

- Do not add steps that were not tested in this session
- Do not present inferences as rules
- High-risk judgments must preserve human confirmation

Third: the compression prompt

If the skill draft is too long for everyday use:

Please condense this skill draft into a version team members can use without reading a manual.

Requirements:

- Lead with the applicable situations
- Keep steps concrete and specific
- Safety constraints and stop conditions must be preserved
- Output format should be directly copyable
- Write it as a workflow, not a tool description

Chapter exercises and self-check:

- Find one workflow you have run two or three times. Evaluate whether it is ready to become a skill.

- Ask AI to debrief that workflow, then judge: is it better as a template, an SOP, or a skill?
- Check: did you avoid packaging a workflow that was not yet stable?

19. Team Adoption and Knowledge Capture

What you will learn:

- Turn personal AI experience into team-shareable templates and processes.
- Build a rhythm for collecting good cases and learning from failures.
- Understand how individual prompts grow into team templates, and team templates grow into skills.

AI assistance becomes reliable only when it moves beyond one person's enthusiasm. The most durable path: each person finds a low-risk scenario, the team distills those into shared templates, and only then does the team consider formalizing skills.

19.1 Individual adoption: start with one small task

Each team member should pick one task that meets all of these conditions:

- Happens at least once a week
- Normally takes 15 or more minutes to organize
- Does not involve formal commitments or high-risk decisions
- Is easy to verify if AI makes an error
- Produces output the next person can act on

Good starting points:

- Meeting note summaries
- Internal summaries of customer messages
- Initial organization of event or project materials
- Expense claim completeness checks
- Open questions extracted from requirement documents

Not suitable as a first AI task:

- Directly replying to a major customer complaint
- Automatically deciding on refunds or compensation
- Automatically modifying official records
- Automatically publishing outward-facing content
- Automatically sending formal correspondence

19.2 Team adoption: collect one good case per week

Use a simple format:

Case name:

Original task: How was this done before? How long did it take?

Assignment method: How did you ask AI to do it?

AI output: What did AI produce?

Human verification: What did you check? Where did AI make errors?

Improvement or time saved: Time saved, fewer missed items, cleaner handoffs, or clearer communication?

What to change next time: How should the prompt or workflow be adjusted?

Do not collect only success stories. Cases where AI made errors have more teaching value - they become verification rules for the next round.

19.3 When a prompt is ready to become a team template

A prompt is ready to become a team template when all of the following are true:

- At least two people have used it
- It has been run at least three times
- Input type is consistent each time
- Output format is consistent each time
- There are clear things AI must not do
- There is a clear way to verify the output

If a prompt has only been used once, keep it as a personal note. Do not rush to make it a template.

19.4 When a team template is ready to become a skill

Chapter 18 covered how an individual creates a skill. For team adoption, the key question is whether a workflow is mature enough to institutionalize.

A team template is ready to become a skill when it is stable enough that:

- Applicable situations are clearly defined
- Input data is predictable
- Operating steps are fixed
- Risk boundaries are explicit
- Stop conditions are clear
- Different team members consistently produce stable results

Do not package everything as a skill. Premature skill-building locks in workflows that have not been fully tested.

19.5 Where to store knowledge

Layer	What it stores	Best for
Personal note	Your own useful prompts and methods	Prompts still being tested or refined
Team template	Reusable formats multiple people use	Meeting summaries, customer message briefs, expense checks
Skill or SOP	Verified, mature workflows	High-risk customer triage, expense reimbursement, email attachment review

Do not store in any shared space:

- Full customer personal data
- Passwords, tokens, or authentication codes
- Unpublished contracts, quotes, or financial data
- Sensitive complete conversations
- Data that only specific people are authorized to view

19.6 Template maintenance rules

Templates go stale without a designated owner. Each team should assign one person to:

- Collect useful new cases
- Remove templates that are no longer applicable
- Add notes about common errors
- Update safety constraints
- Confirm that no template encourages AI to make decisions beyond its authorized scope

When updating a template, record why it changed. For example:

2026-05-05 update:
In several customer service cases, AI was prone to directly promising refunds.
Added constraint: "Do not commit to refunds, compensation, or resolution timelines."

This kind of note helps the next person understand which edge cases the template has already accounted for.

19.7 How managers can drive adoption effectively

Do not begin by asking everyone to fully adopt AI. A more workable pace:

1. Each person picks one low-risk task to try first.
2. Share one success or failure case per week.
3. Collect recurring good prompts into team templates.
4. Add human approval checkpoints for high-risk tasks.
5. After one month, evaluate which workflows are ready for skill packaging.

What managers should measure is not how often people used AI. The right questions are:

- Does it save organizing time?
- Does it reduce missed items?
- Does it make handoffs cleaner?
- Does it lower the risk of accidental commitments?
- Does it produce reusable templates?

19.8 Three small things to try first

There is no need to start with the biggest task. Try one of these as a first run:

1. Take an application form, event brief, or requirements document and ask AI to extract the key points, risks, and next steps.
2. Take a system-exported report and ask AI to spot anomalies, clean up columns, and summarize in plain language.
3. Take a manager's task handoff or a set of meeting notes and ask AI to break it into steps, list open questions, and produce a first draft.

After trying, do not just ask whether it was useful. Ask these instead:

- Will this task recur in the future?
- Did the AI first draft actually save time?
- Was the output easy to verify?
- What constraints should be stated upfront next time?
- Is this prompt worth saving?

If the answers are mostly yes, save it as a personal note. Once multiple people have used it, it has run several times, and verification is stable - then move it to a team template or skill.

19.9 One person's skill is personal speed. A shared team library is organizational speed.

If everyone keeps their best methods inside their own conversations, the team accumulates nothing.

The real value comes from moving good workflows into a shared space such as a team knowledge base, a shared wiki, a shared folder, an internal SOP document, or a shared AI skill library.

When sharing, always include four things:

- What situation this skill or template applies to
- What data to prepare before using it
- How to verify the output
- What situations require human judgment instead of AI

Do not just paste a raw prompt. A prompt without context, boundaries, and a verification method is easy to misuse.

Chapter 19 exercises:

- Collect one success case and one case where AI made an error.
- Decide whether each belongs in a personal note or a team template.
- Check: does your team have a shared storage location and a designated maintainer?

20. One-Page Quick Reference

What you will learn: Keep this chapter accessible during daily work. It gives you the key questions, templates, and decision guides without re-reading the full handbook.

20.1 Before assigning

Ask yourself:

- What do I want AI to complete?
- What data can I provide?
- What must it not do?
- What format should the output be?
- How will I know if the output is usable?

20.2 When reviewing output

Check:

- Did it miss anything?
- Did it guess something not in the source?
- Did it overstate or imply commitments?
- Does it violate any company policy?
- Did it expose sensitive information?
- Can the next person act on this directly?

20.3 Universal quick template

Please help me handle the following.

Goal:

[What you need]

Data:

[Paste content here]

Constraints:

- Do not add information not present in the source
- Label anything uncertain as "to be confirmed"
- Do not make final decisions for me
- Do not add commitments that are not in the source

Output:

1. Summary
2. Confirmed items
3. Items to confirm
4. Risks
5. Suggested next step

20.4 High-risk task quick check

When the task involves refunds, compensation, quotes, discounts, contracts, legal matters, tax, HR, external announcements, formal correspondence, customer commitments, deletion, overwriting, publishing, permission changes, personal data, financial transactions, or confidential information:

AI can only help organize or draft. Use this safe assignment phrase:

Please organize the facts, inferences, open questions, and risks. Do not make final decisions for me and do not add commitments.

20.5 Key principles

- Chatting asks for an answer. Delegation starts work.
- AI's first version is a draft, not a final answer.
- AI can speed you up, but you set and guard the boundaries.
- You do not need many AI tools. Use a few mainstream ones well.
- The problem is often not that the tool cannot do it - it is that you have not learned how to delegate it clearly.
- Do not fear AI entering your work. Fear handing over your judgment along with the task.
- A skill is not a saved prompt. It is a verified, repeatable workflow.
- One person using AI well speeds themselves up. A team sharing templates and skills speeds the whole organization up.

20.6 How to evaluate a new AI tool

When someone recommends a new AI tool, ask these questions first:

- Can ChatGPT, Claude, or Gemini already handle this with a better-structured prompt?
- Is the real gap that I have not clearly defined the goal, data, constraints, and output format yet?
- Does this tool require logging into company systems, reading email, connecting to data sources, or gaining special permissions?
- Will it touch personal data, financial transactions, contracts, customer records, or confidential information?
- If a mainstream AI combined with a good template can handle it, is an additional tool actually necessary?

A simple guideline:

For general summarizing, rewriting, categorizing, breaking down, or checking - use a mainstream AI first. For system integration, batch automation, specialized formats, or formal regulated workflows - then evaluate a dedicated tool.

20.7 Chat / workspace agent / skill decision table

What you are doing now	Recommended approach	Reason
Asking a concept, exploring an idea, rewriting a short passage	General Chat	No file access needed; conversation context is sufficient
Summarizing pasted content	General Chat	Small data; works within conversation context
Drafting a prompt or skill	General Chat or workspace agent	Think it through first before writing to disk
Reading multiple files, documents, or a folder	Codex / Claude Code / Gemini CLI	Requires workspace context beyond what you can paste
Modifying files, organizing folders, running checks	Codex / Claude Code / Gemini CLI	Requires actual read-write and verification
Writing a verified workflow as a formal skill file	Workspace agent or manually	Needs confirmed paths, format, and file structure
A workflow multiple people use with consistent format	Template first, skill once it is stable	Confirm the workflow is reliable before formalizing
Deletion, sending, publishing, permission changes, financial transactions, or personal data	Stop and confirm with a human	This is not a tool choice question - it is a responsibility and safety boundary

In one sentence:

Chat is for thinking clearly and producing drafts. Workspace agents put things into action. Skills preserve what is already verified, repeatable, and worth reusing.

Chapter 20 exercises:

- Use this quick reference to re-evaluate one real task you have done recently.
- Decide whether to use Chat, a workspace agent, a template, or a skill.
- Check: can you find the safe assignment phrase within one minute?

21. Closing: From Tool User to AI Workflow Builder

What you will learn: Consolidate the core ideas from this handbook. Understand the shift from using AI as a chat tool to directing it as a reliable work assistant.

AI does not give a perfect answer the first time.

A more realistic expectation:

1. AI produces a first draft.
2. You find the problems.
3. You specify what needs to change.
4. AI produces a revised draft.
5. You save the useful workflow for next time.

Using AI well is not about memorizing commands.

What actually matters:

- Knowing what to delegate
- Knowing what not to delegate
- Seeing where AI went wrong
- Explaining clearly how to fix it
- Turning good methods into repeatable workflows

Start with low-risk, small tasks. When you can reliably complete the assign-verify-refine cycle, an AI work assistant stops being a chat tool and starts being a genuine work partner.

The concern that AI will replace workers is real. The practical answer is not to try to generate text faster than AI can. It is to become someone who can define work clearly, delegate to AI precisely, verify outputs rigorously, and improve processes systematically. People who can do all four of those things - define, delegate, verify, improve - are harder to replace than those who only produce output. They become the ones who build methods that others in the team can rely on.

The person who can direct AI to produce verifiable, deliverable, and continuously improving results is not just a tool user. They are a workflow builder, a verifier, and a process designer. Their value in a team is not "I completed a task." It is "I built a workflow that the whole team can reuse."

Chapter 21 exercises:

- Write down the first low-risk task you will delegate to AI in the coming week.
- Decide where to save the result: personal note, team template, or shared library.
- Check: have you shifted from asking for answers to building repeatable workflows?

Appendix A: Skill Example - High-Risk Customer Message Triage

This skill is not for labeling customers. It helps customer-facing staff organize facts, reduce the risk of accidental commitments, and determine whether a case needs escalation — especially in high-emotion, high-pressure situations. Use the term "high-tension customer communication" rather than labels that could demean the person.

A.1 Skill Name

High-Risk Customer Message Triage

A.2 Applicable Situations

Use when a customer service interaction involves:

- Strong customer emotion or repeated accusations
- Requests for refunds, compensation, discounts, or special exceptions
- Threats to complain, post negative reviews, or take legal action
- Repeated topic-jumping or demands that exceed policy
- Denial of previous conversation records
- Demands for an immediate commitment from customer service
- Threatening, humiliating, or emotionally coercive language
- Risk of escalation to PR, legal counsel, or senior management

A.3 Required Data

1. **Customer's original message** — LINE, email, form submission, or call summary. Preserve exact wording; do not rely on memory or paraphrase.
2. **Case background** — order, service, event, contract, and payment status; timeline of what happened.
3. **Previous replies and commitments** — what customer service has already said, whether a refund, compensation, or deadline has been promised.
4. **Company rules and constraints** — refund policy, service terms, compensation standards, what customer service is authorized to do.
5. **What customer service wants to achieve** — de-escalate, clarify facts, decline an unreasonable request, escalate to management, preserve a paper trail, or prepare a formal reply.

A.4 Assignment Template

Please help me triage the following customer case using the High-Risk Customer Message Triage method.

Goal:

Do not reply to the customer yet. First help me assess the case status, risks, possible response directions, and whether escalation is needed.

Data:

[Customer's original message]

[Case background]

[Previous commitments or replies from our side]

[Company rules or constraints]

Constraints:

- Do not commit to refunds, compensation, discounts, or exceptions on behalf of the company
- Do not use language that demeans the customer
- Do not state inferences as facts
- Do not suggest customer service argue with the customer
- Do not provide legal conclusions
- Label anything uncertain as "to be confirmed"
- If legal, PR, data privacy, or large financial issues are involved, flag as "escalation required"

Output:

1. Case summary
2. Customer's main request
3. Customer tone and communication risk
4. Confirmed facts
5. Unconfirmed items
6. What our side has committed to
7. What we cannot commit to
8. Risk level
9. Recommended handling strategy
10. Suggested reply direction
11. Whether escalation to manager / legal / PR is needed
12. Internal case notes

A.5 Risk Levels

L1 — General dissatisfaction: Customer is unhappy but still describing the problem. No threats of complaint, negative review, or legal action. Customer service can respond per standard procedures.

L2 — High-emotion complaint: Strong customer emotion or repeated demands for exceptions. Requires a more careful tone, but still manageable by customer service without escalation.

L3 — Escalation risk: Customer threatens to complain, post a negative review, or make the situation public. The case involves a refund dispute, compensation claim, personal data issue, or payment discrepancy. Manager should confirm the response before it is sent.

L4 — Legal or PR risk: Customer mentions lawsuits, lawyers, regulators, or media. The case involves a major financial dispute, data breach, or contract interpretation. Customer service should not respond independently — hand off to legal, PR, or senior management.

A.6 Recommended Handling Strategy

For L1:

- Confirm what the problem is
- Provide a clear handling step
- Keep the tone steady and the message concise

For L2:

- Acknowledge the customer's emotion without committing to an outcome
- Restate the confirmed facts
- Lay out a clear next step
- Avoid any wording that reads as a counterargument

For L3:

- Pause all commitments until management confirms what can be said
- Gather the complete case file
- Reply to the customer with a holding message: "We have received your message and are following up."

For L4:

- Do not make any legal liability judgment
- Do not admit fault
- Do not promise compensation
- Do not delete any records
- Preserve the full conversation and handling timeline
- Transfer the case to management, legal, or PR

A.7 Reply Direction

Language to use:

- "We have received your feedback."
- "We will review the relevant records before responding."
- "We need to clarify [specific item] first."
- "Once confirmed, we will reply with the next steps available to us."
- "This requires confirmation through our internal process. We cannot make a direct commitment at this stage."

Language to avoid:

- "You misunderstood."
- "This is not our problem."
- "Rules are rules."
- "That's impossible."
- "You said that yourself before."
- "We will definitely give you a refund."

- "We guarantee this will be resolved."

A.8 Draft Reply Templates

High emotion, not yet escalated:

Thank you for reaching out. We understand you are frustrated with what happened, and we take that seriously.

To make sure we have the full picture, we will review the relevant records and the current status of your case before responding further. Specifically, we need to clarify:

1. [item]
2. [item]

Once we have confirmed those details, we will follow up with the next steps we can take. Until then, our team will not be making any commitments regarding refunds, compensation, or exceptions, so we can ensure the information we provide is accurate.

Customer demands an immediate refund:

Thank you for your message. We have received your refund request.

Whether this refund can be processed depends on the order status and our company policy. We need to verify:

1. Order status
2. Payment status
3. Whether the service has been fully delivered
4. Whether there is any prior handling record on file

We will follow up with the available options after confirming. Our team cannot commit to a refund outcome before completing this review, and we appreciate your understanding.

Customer threatens a negative review or formal complaint:

Thank you for letting us know. We take your situation seriously.

To respond accurately, we will compile the relevant records – including your message, order details, and our previous communications – and then reply in accordance with our internal process.

Until we have completed that review, our team will not be making any judgment on refunds, compensation, or responsibility, to avoid any confusion from incomplete information.

A.9 Stop Conditions

AI should flag "escalation required" and not suggest customer service handle the case independently when:

- Customer mentions legal counsel, a lawsuit, or a regulatory authority
- Customer threatens media coverage or public disclosure
- A personal data breach is suspected
- Large financial amounts or formal compensation are involved
- Contract terms require interpretation
- Physical threats or harassment are present

- Customer demands deletion of records
- Customer requests internal company information
- Customer demands that company procedures be bypassed
- Customer service has responded multiple times but the conflict continues to escalate

A.10 Verification Checklist

After using this skill, verify:

- Facts and inferences are clearly separated
- Committed and uncommitted items are listed
- No language that could escalate the situation or demean the customer
- No direct commitments to refund, compensation, or discount
- Escalation need has been assessed
- Case timeline and conversation records are preserved
- Draft reply can be reviewed by a manager

Appendix B: Skill Example - Expense Reimbursement Completeness Check

This skill is not for AI to decide whether an expense can be reimbursed. It helps team members organize data, check for missing items, and write clear purpose descriptions before submitting to finance. Finance review and final approval always follow company policy.

B.1 Skill Name

Expense Reimbursement Completeness Check

B.2 Applicable Situations

Use when a team member needs to submit an expense claim and wants to:

- Verify that all required documents are present
- Identify missing items before sending to finance
- Write a clear purpose description
- Avoid back-and-forth follow-up requests from the finance team

Applicable to: petty cash, transportation, purchases, event costs, business meals, software subscriptions, travel and accommodation, advance payment requests, supplemental receipts.

B.3 Required Data

Basic expense data:

- Expense item
- Date incurred
- Amount and currency
- Payment method
- Payer or advance payer
- Department or project
- Purpose description

Supporting documents:

- Invoice or receipt
- Electronic invoice proof
- Credit card statement
- Bank transfer screenshot
- Order or booking screenshot

- Contract or quote
- Event or travel approval record

Company rules:

- Reimbursement deadline
- Document format requirements (company name, tax ID on invoice)
- Approved and excluded expense types
- Manager approval threshold
- Whether pre-approval is required for certain categories

B.4 Assignment Template

Please help me check this expense reimbursement using the Expense Reimbursement Completeness Check method.

Goal:

Check whether the data is complete, identify anything missing, and flag anything that needs a clearer explanation. Do not decide whether the expense can be approved. Do not make the final call for finance.

Data:

[Expense item]
[Date]
[Amount and currency]
[Payment method]
[Payer]
[Department or project]
[Purpose]
[Documents provided]
[Company reimbursement rules]

Constraints:

- Do not make the final approval decision
- Do not promise the expense can be reimbursed
- Do not infer documents that were not provided
- Do not ask for full credit card numbers, bank accounts, or ID numbers
- Flag sensitive fields (invoice number, card number, account number) for masking
- Label anything insufficient as "to be supplemented / to be confirmed"

Output:

1. Reimbursement summary
2. Data provided
3. Missing items list (required / recommended / uncertain)
4. Items that need a clearer explanation
5. Items that may need manager approval
6. Pre-submission checklist
7. Suggested purpose statement text
8. Items that finance needs to confirm

B.5 Output Format Example

1. Reimbursement summary

Field	Value
Expense type	

Field	Value
Date	
Amount	
Payer	
Project	
Purpose	
Status	Complete / Missing items / Needs manager confirmation / Needs finance judgment

2. Documents provided

Document	Status
Invoice or receipt	Present / Missing / Unclear
Payment proof	Present / Missing / Unclear
Purpose statement	Present / Missing / Insufficient
Project or department assignment	Present / Missing
Manager approval record	Present / Missing / May be needed

3. Missing items

Required — submission will likely be rejected without these:

- (list each required missing item)

Recommended — reduces follow-up risk:

- (list each recommended addition)

Uncertain — confirm with finance:

- (list items whose necessity is unclear)

Common missing items:

- Invoice or receipt not attached
- Payment proof not attached
- Purpose description absent or vague
- Project or department not specified
- Manager approval record not included
- Company name or tax ID on invoice does not match
- Amount differs from payment screenshot
- Date is past the reimbursement deadline

B.6 Suggested Purpose Statement Templates

General purchase:

This expense covers [item] purchased for the [project name] project, used for [specific purpose]. Advance paid by [name]. Invoice and payment proof attached.

Transportation:

This is a transportation expense incurred for a [business meeting / event / site visit] on [date]. Route: [origin] to [destination]. Relevant boarding pass or payment proof attached.

Event or activity cost:

This expense was incurred for [event name] on [date], covering [cost items]. Purpose and event approval record attached.

Software or tool subscription:

This is a subscription fee for [tool / service], used for [purpose]. Subscription period: [start date] to [end date]. Payment confirmation and subscription record attached.

Business meal or client entertainment:

This expense covers a [client meeting / team working meal] on [date], attended by [names or roles], for the purpose of [meeting objective]. Receipt attached. Please confirm whether pre-approval is required for this category.

B.7 Pre-Submission Checklist

Before sending to finance, verify:

- ☐ Dates on all documents are correct
- ☐ Amount matches the supporting documents
- ☐ Currency is correct
- ☐ Invoice or receipt is legible
- ☐ Company name and tax ID on the invoice match company requirements
- ☐ Payment proof is attached
- ☐ Purpose description is clear
- ☐ Department or project is specified
- ☐ Manager approval obtained if required
- ☐ Submission is within the reimbursement deadline
- ☐ No personal expenses are mixed in
- ☐ Sensitive fields (card number, account number) are masked
- ☐ All attachments are uploaded

B.8 Sensitive Data Reminders

Do not paste the following into a conversation or shared document:

- Full credit card numbers
- Bank account numbers
- National ID numbers
- Home addresses
- Complete salary information
- Unmasked company account information
- Payment platform login credentials
- API keys or access tokens

Safer alternatives:

- Use only the last four digits of card or account numbers
- Refer to clients or vendors as "Client A" or "Vendor B"
- Use partial order numbers
- Mask sensitive fields in screenshots before sharing

B.9 Stop Conditions

AI should stop and flag for finance or manager review when:

- Documents appear incomplete or internally inconsistent
- Amount does not match the supporting document
- Company name or tax ID is incorrect
- Submission deadline has passed
- Personal and business expenses appear mixed
- Large amounts are involved
- Entertainment, gifts, commissions, or consulting fees are included
- Foreign currency or overseas payment is involved
- Contract, procurement, or subscription renewal is involved
- Tax treatment questions arise
- Supplemental invoices or document amendments are needed
- Purchase was not pre-approved

B.10 Verification Checklist

After using this skill, verify:

- Provided documents are clearly listed
- Missing items are clearly identified
- Finance has not been replaced as the final approver
- No commitment has been made that the expense will be approved
- Sensitive data fields are flagged for masking
- Items needing manager or finance confirmation are labeled
- Purpose description matches the actual use
- Output is useful as a pre-submission self-check

Appendix C: Skill Example - Email Search and Attachment Safety Check

This skill helps users organize email search conditions, identify candidate messages, check attachment risks, and produce a delivery checklist. The public edition does not recommend letting AI directly operate a private inbox, download attachments, or extract full message content. **AI can only assist with search conditions and operating steps unless the organization has authorized tools or connectors and the user has given explicit permission.** If no authorization exists, AI must be transparent about that limitation and must not pretend to have searched or found anything.

C.1 Skill Name

Email Search and Attachment Safety Check

C.2 Applicable Situations

Use when you know a specific email exists and need help with one or more of the following:

- Organizing search conditions so you can find it yourself
- Deciding which keywords or filters to use
- Comparing candidate messages against your criteria
- Checking attachment type and risk before downloading
- Preparing a delivery checklist before submitting screenshots or attachments as supporting evidence

Common use cases: e-invoice emails, receipt emails, platform order confirmations, payment notifications, vendor invoices, event registration confirmations, travel booking confirmations, software subscription receipts.

C.3 Required Information

Provide as many of the following as you have:

- Sender (email address or recognizable name)
- Subject keywords
- Date range
- Amount, order number, or other unique identifiers
- Body content keywords
- Whether an attachment is expected

C.4 Assignment Template

Please help me organize search conditions and a pre-delivery checklist for the following email.

Search clues:

- Sender:
- Subject keywords:
- Date range:
- Amount or order number:
- Content keywords:

I need:

1. Suggested search conditions
2. Candidate email matching criteria
3. What to mask before screenshotting the email
4. What attachment risks to check before downloading
5. If my email tool is authorized and capable, list the deliverable files

Constraints:

- Do not read or report unrelated email content
- Do not download suspicious attachments
- Do not click unknown links
- Do not output full credit card numbers, bank accounts, verification codes, or passwords
- If you do not have actual inbox access, clearly state you can only provide search conditions and checking steps – do not pretend to have searched or found anything
- If multiple similar emails are found, list candidates for me to confirm; do not choose for me

C.5 Operation Flow

Step 1 — Define search conditions

Using the clues provided, suggest:

- Sender filter
- Subject keywords
- Date range
- Amount, order number, or body content keywords
- Whether to filter for emails with attachments only

Even when only one result appears, ask the user to confirm before proceeding. The public edition discourages AI from opening large numbers of unrelated messages.

If multiple candidates appear, list them:

Found 3 possible matches:

1. 2026-05-01 | Apple | Your receipt from Apple | \$9.99 | Has attachment
2. 2026-05-02 | Apple | Your receipt from Apple | \$29.99 | Has attachment
3. 2026-05-03 | Apple | Subscription renewal | \$9.99 | No attachment

Please confirm which one to proceed with.

Step 2 — Confirm email content

Once a match is selected, confirm:

- Sender
- Date
- Subject

- Whether an attachment is present
- Whether it matches the amount, order number, or keywords provided

Do not paste the full email body into the conversation unless the user specifically requests it.

Step 3 — Before taking a screenshot

A useful screenshot for reimbursement or evidence purposes should include:

- Sender
- Date
- Subject
- Key content (amount, item description, order number)
- Attachment section (if any)

Before screenshotting, check:

- Are full credit card numbers, phone numbers, home addresses, or verification codes visible? Mask or crop them before sharing.
- If no masking capability is available, advise the user which fields to cover manually.

Step 4 — Before downloading an attachment

Common safe attachment types: PDF, JPG, PNG, CSV, XLSX, receipts, invoices, order confirmations.

Before downloading, verify:

- Attachment comes from a trusted sender
- File extension matches what is expected
- It is not a .exe, .dmg, unknown .zip, or macro-enabled file
- It does not require logging into an external website to access
- It does not appear to contain full financial account data or government-issued ID information

Step 5 — Verify delivered files

If files were actually downloaded or screenshots taken, verify:

- The file exists and is not 0 bytes
- The file extension matches expectations
- The PDF or image can be opened
- The number of attachments matches what the email showed

C.6 Output Format

Search and match results:

- Suggested search conditions:
- Suggested date range:
- Suggested keywords:
- Candidate matching criteria:

- Fields that need user confirmation:

Pre-delivery safety check:

- What to mask in the screenshot:
- What to confirm before downloading the attachment:
- What not to download or click:

If the tool is authorized and processing has occurred:

- Email confirmed:
- Screenshot file:
- Attachment file:

Outstanding / notes:

-

C.7 When the Email Cannot Be Found

No matching email found with the current search conditions.

Searched:

- Sender:
- Subject:
- Date range:
- Keywords:

Possible reasons:

- Date range is too narrow
- Sender name varies by platform or region
- Email is in spam, archive, or a subfolder
- Subject line uses different wording than expected
- Attachment was delivered via a link rather than as a file

Suggested next steps:

- Expand the date range
- Try searching by amount or order number instead
- Search for: invoice / receipt / order confirmation / payment successful

C.8 Security Limits

AI must not:

- Request the user's email password or any verification code
- Click suspicious payment links
- Download .exe, .dmg, or unknown .zip files
- Open or read unrelated emails
- Paste full sensitive information (financial data, ID numbers) into the conversation
- Forward emails to any third party
- Delete, archive, move, or mark emails as read or unread without explicit user instruction

C.9 Stop Conditions

Stop and ask the user before continuing when:

- Multiple highly similar candidates exist and the user has not confirmed which one to use

- The email contains sensitive personal data or full financial account information
- The attachment type is unexpected or suspicious
- The email looks like phishing
- Downloading the attachment requires logging into an external site
- Any forwarding, deleting, archiving, or moving of emails is involved
- Email content does not match the user's description

C.10 Verification Checklist

After completing the task, verify:

- The email found matches the user's specified clues
- If a screenshot was taken, sensitive fields are masked
- If an attachment was downloaded, the type, source, and file integrity have been confirmed
- No unrelated or suspicious files were downloaded
- No sensitive information was output to the conversation
- Search conditions, candidate criteria, and file list have been documented
- If multiple candidates existed, the user confirmed the correct one — no guessing

Appendix D: Skill Example - Safe Computer Storage Cleanup

The focus of this skill is safe read-only surveying and risk classification — not encouraging users to let AI delete files. Before any actual cleanup, back up first and confirm each item individually. Stop immediately if company devices, cloud-synced data, customer records, financial data, contracts, passwords, or system configuration are involved.

D.1 Skill Name

Safe Computer Storage Cleanup

D.2 Applicable Situations

Use when your computer is low on storage and you want to:

- Find what is using the most space
- Identify cleanup candidates safely, without installing unknown software
- Know what is safe to delete and what must never be touched

Not suitable for:

- Suspected malware or account compromise (contact IT security instead)
- Clearing traces of confidential or privacy-sensitive data
- Reinstalling the operating system
- Bypassing company device management or access controls
- Asking AI to delete large amounts of data without item-by-item confirmation

D.3 The Real Risk with "Free Cleanup Software"

Many tools that claim to solve low-disk problems ask you to: download a .dmg, install a profile, enter your system password, and grant Full Disk Access. Once granted, such tools can silently read your browser sessions, email, cloud storage, and saved credentials. The worst outcome is not a messy hard drive — it is losing account access, having cloud data compromised, or exposing work credentials.

Use a workspace AI agent to do a safe, read-only survey first. Do not install unknown tools.

D.4 Information to Provide Before Starting

Provide:

- Operating system (e.g., macOS or Windows)
- Current symptom (e.g., cannot update, cannot download, system is slow)
- Whether the device is company-issued or managed by IT
- Any folders or projects that must not be touched

- Whether a Time Machine backup, external drive, or cloud sync is in place

If this information is not provided, the assistant should ask first:

Before we begin, please confirm:

1. Is this a company-issued or IT-managed device?
2. Are there company backup or data-retention rules that apply?
3. Which folders are completely off-limits?
4. Do you have a Time Machine, cloud, or external drive backup in place?
5. Would you like to start with a read-only survey only, or can I proceed to move confirmed items to the trash after you review each one?

D.5 Assignment Template

Please help me do a safe cleanup of low disk space.

Goal:

Identify files and folders that can be safely cleaned up to free storage space, without affecting daily use, work projects, account security, or system stability.

Constraints:

- Do not delete any files yet
- Do not download or install any external cleanup software
- Do not touch passwords, browser data, email, cloud-synced folders, work projects, or system settings
- Do not use permanent deletion (bypassing the trash)
- Do not assume "large file" or "not recently modified" means "safe to delete"
- Start with a read-only survey

Please output:

1. Current disk space remaining
2. Largest folders or file types by size
3. Low-risk candidates recommended for review
4. Candidates that need my confirmation before touching
5. High-risk items that must not be touched
6. Estimated space that could be freed

Format as a table: path, size, type, likely use, frequency (frequent / infrequent / unknown), suggested action, deletion risk, whether my confirmation is needed.

D.6 Operation Flow

1. Confirm the operating system, whether the device is company-managed, and backup status.
2. Explicitly remind the user not to download unknown cleanup software.
3. Run a read-only survey of remaining space and the largest folders.
4. Compile a candidate list — no deletions at this stage.
5. Classify candidates into: low-risk, needs confirmation, high-risk do-not-touch.
6. Explain the likely purpose and deletion impact for each candidate.
7. Only after the user explicitly confirms each item, list the handling steps for low-risk items.
8. If the user proceeds, move items to the trash — do not permanently delete.
9. After cleanup, re-survey disk space to verify the change.

10. Report what was processed, how much space was freed, what was skipped and why, and any follow-up recommendations.

D.7 Classification Rules

Low-risk candidates (still need user confirmation before touching):

- Old installers, archives, and duplicate exports in the Downloads folder
- Screen recordings, meeting recordings, or temporary transcoded files no longer needed
- Caches and build artifacts that can be regenerated and do not affect login, sync, or project state
- Installer packages or exported reports that can be re-downloaded
- Temporary project assets that have been delivered, backed up, and confirmed no longer needed

Candidates that need confirmation:

- Large video, image, or design asset files
- Old project folders
- Development build artifacts, simulator data, package manager caches
- Local copies of cloud-synced folders
- Compressed archives whose delivery or backup status is uncertain

High-risk items — AI must not delete these without explicit, item-by-item user confirmation:

- Passwords, Keychain entries, certificates, tokens, API keys
- Browser profiles, cookies, sessions, saved login data
- Email, Messages, and communication history
- Work projects, customer data, financial records, contracts, original expense documents
- iCloud Drive, Google Drive, Dropbox, and other cloud-synced folders
- Application Support folders for apps whose function is unclear
- System folders, permission settings, and device management configurations
- Photo libraries, music libraries, and video editing project source files

D.8 Output Format

Survey output:

Current status:

- Total disk capacity:
- Used:
- Remaining:
- Recommended minimum free space to maintain:

Top space users:

Rank	Path / Type	Size	Initial assessment	Notes
1				
2				

| 3 | | | |

Cleanup candidates:

Path	Size	Type	Frequency	Suggested action	Risk	Needs confirmation?
----	----	----	-----	-----	----	-----
						Yes / No

Items not to touch:

- [path]: reason
- [path]: reason

Suggested next steps:

1. Review low-risk candidates first
2. Back up and confirm the purpose of each item before acting
3. Move to trash only – do not permanently delete
4. Wait 3–7 days to confirm no impact before emptying the trash

Post-cleanup report:

Cleanup results:

- Disk space before cleanup:
- Disk space after cleanup:
- Space freed:

Moved to trash:

Path	Size	Type	Confirmed by
----	----	----	-----

Skipped items:

Path	Reason
----	-----

Follow-up reminders:

- Use the computer normally for a few days before emptying the trash
- If storage is still insufficient, schedule a second-round survey
- If a system software update triggered the low-space warning, the system partition itself may be the primary consumer

D.9 Security Limits

AI must not:

- Download or install any disk cleanup tool
- Ask for the user's system password, Apple ID, Google account, or company credentials
- Request Full Disk Access permissions unless an authorized IT process specifically requires it
- Use permanent deletion (skip the trash)
- Execute any deletions before backup and item-by-item user confirmation
- Delete any path not in the confirmed candidate list
- Independently clean browser data, email, passwords, cloud-synced folders, or system configuration
- Share full file paths containing personal names or private identifiers with third parties
- Use "not recently modified" as the sole reason to recommend deletion
- Overstate cleanup results as a guaranteed fix

D.10 Stop Conditions

Stop and ask the user before continuing when:

- User requests downloading unknown cleanup software
- User asks for bulk deletion without item-by-item confirmation
- System password or elevated permissions are required
- Candidates include company data, customer data, financial records, contracts, or personal information
- Candidates are inside a cloud-synced folder
- A folder's purpose is unclear but it is large
- Survey results before and after cleanup are inconsistent
- User has no backup but wants to delete a folder that may contain important data
- The device shows signs of malware or account compromise

D.11 Verification Checklist

After completing the task, verify:

- AI started with a read-only survey — no deletions at the outset
- No external cleanup software was downloaded or installed
- Cleanup candidates and their risk levels were listed
- Items were classified as frequent / infrequent / unknown
- Browser data, email, passwords, cloud-synced folders, company data, and system settings were not touched
- Only user-confirmed items were processed
- Items were moved to the trash, not permanently deleted
- Disk space before and after cleanup was reported

AI Work Assistant Handbook - Free Public Edition v1.0